

Complex Libraries

Using Hash Dictionaries

Playing Hash Table

*You are the new produce manager of the local grocery store.
You want to use a dictionary to track your fruit inventory.*

Entries have the form

(“banana”, 20)

where

- “banana” is the **key**
 - 20 is the associated data, like the number of cases in stock
-
- Let’s observe your initial interactions with a hypothetical hash dictionary library

This is your side

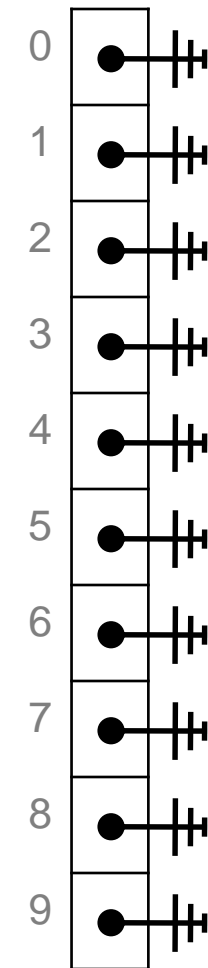
This is what
is going on in
the library

Client | **Implementation**

*You begin by
creating a
new dictionary*

Create a new hash dictionary

Here you go!



- This library uses separate-chaining hash tables to implement dictionaries
- It decides on an initial capacity of 10
 - it's probably self-resizing

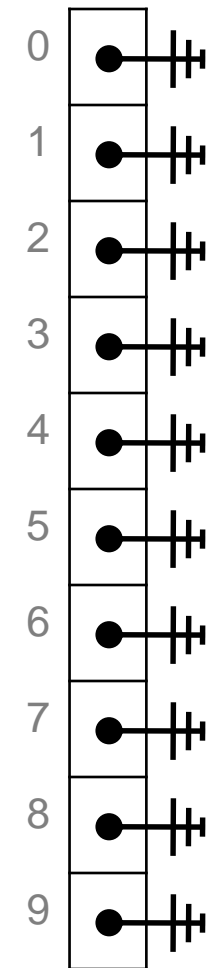
Client | Implementation

Insert A = ("apple", 20)

What's the key of (A)?

✓ new dictionary

Next, you *insert*
A = ("apple", 20)



- Why is the library asking this?
 - it does not know what entries are
 - (A) is just a pointer to some struct
 - no sense of what's in it
- ***You need to tell it***

Client Implementation

Insert A = ("apple", 20)

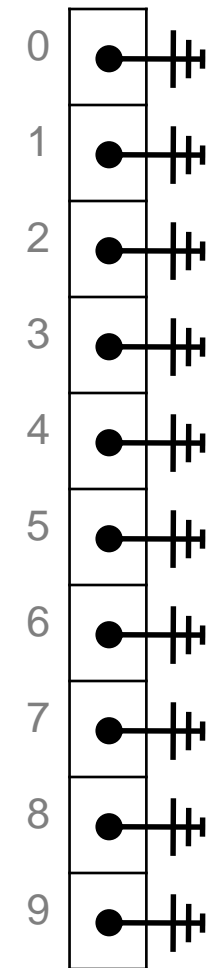
What's the key of (A)?

"apple"

What's its hash value?

✓ new dictionary

Next, you *insert*
A = ("apple", 20)



- Why is the library asking this?
 - it does not know the type of keys
 - even if it did, there are many ways to hash them
- ***You need to tell it***

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

✓ new dictionary

Next, you *insert*
A = ("apple", 20)

Client Implementation

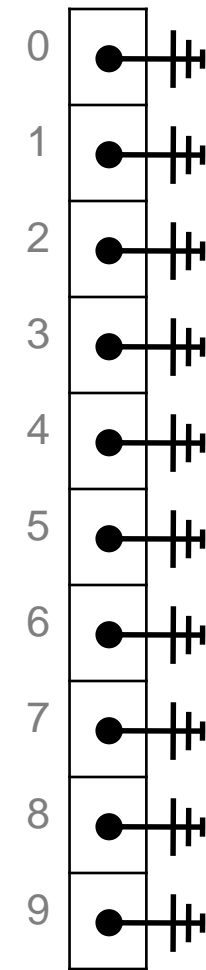
Insert A = ("apple", 20)

What's the key of (A)?

"apple"

What's its hash value?

-1290151091



- $-1290151091 \% 10$ is -1 in C0
 - not a valid array index! ✗
 - the library needs a more robust way to compute the hash index
- Let's say it keeps the last digit

Exercise!

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

✓ new dictionary

Next, you *insert*
A = ("apple", 20)

Client Implementation

Insert A = ("apple", 20)

What's the key of (A)?

"apple"

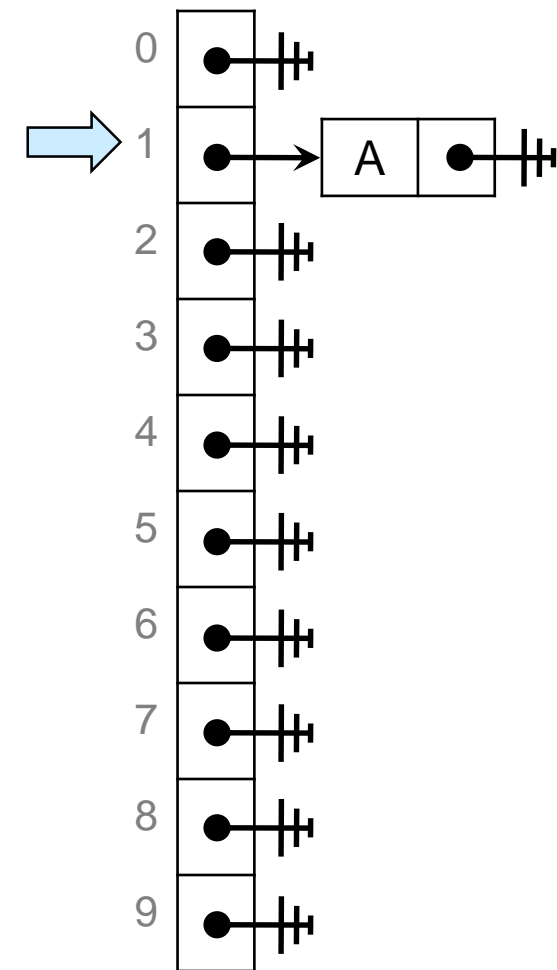
What's its hash value?

-1290151091

Ok. The
hash index is 1.

This chain is empty.
I can insert entry
(A) there.

Done



- The library asked for
 - the key of the entry
 - the hash value of the key

Funny! Libraries didn't ask
for anything in the past

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

Client Implementation

Insert B = ("banana", 10)

What's the key of (B)?

"banana"

What's its hash value?

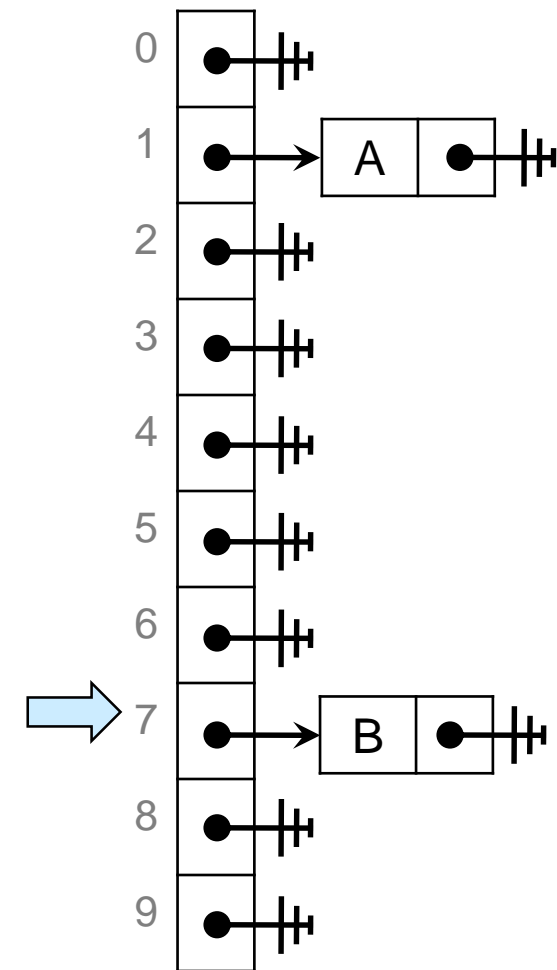
207055587

Ok. The
hash index is **7**.

This chain is empty.
I can insert entry
(B) there.

Done

Same as for (A)



- ✓ new dictionary
- ✓ insert A = ("apple", 20)

Next, you *insert*
B = ("banana", 10)

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

Client Implementation

Insert C = ("pumpkin", 50)

What's the key of (C)?

"pumpkin"

What's its hash value?

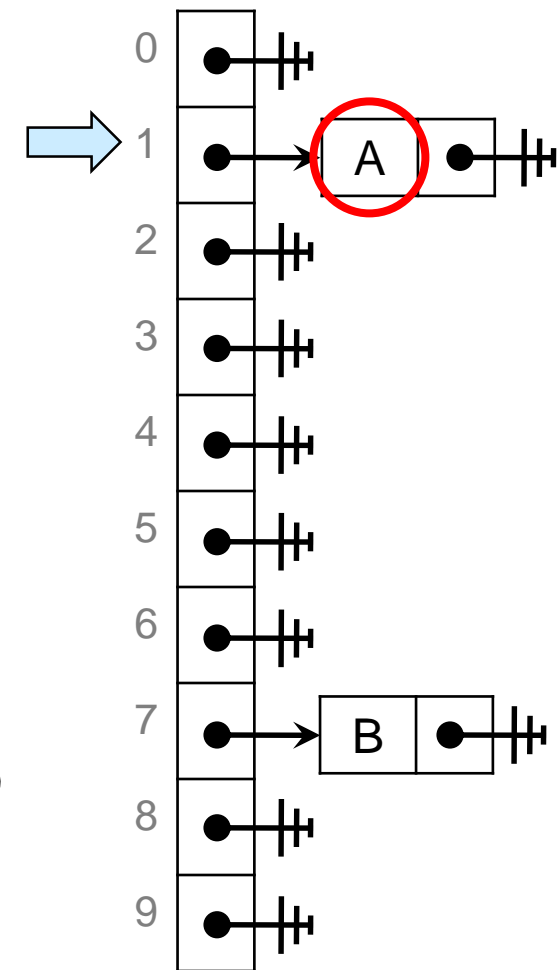
-1189657311

Ok. The hash index is 1.
It points to a node
for entry (A)

What's the key of (A)?

- ✓ new dictionary
- ✓ insert A = ("apple", 20)
- ✓ insert B = ("banana", 10)

Next, you *insert*
C = ("pumpkin", 50)



- Why is the library asking this?
 - it does not know what entries are
 - (A) is just a pointer to some struct
 - no sense of what's in it
- ***You need to tell it***

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

Client Implementation

Insert C = ("pumpkin", 50)

What's the key of (C)?

"pumpkin"

What's its hash value?

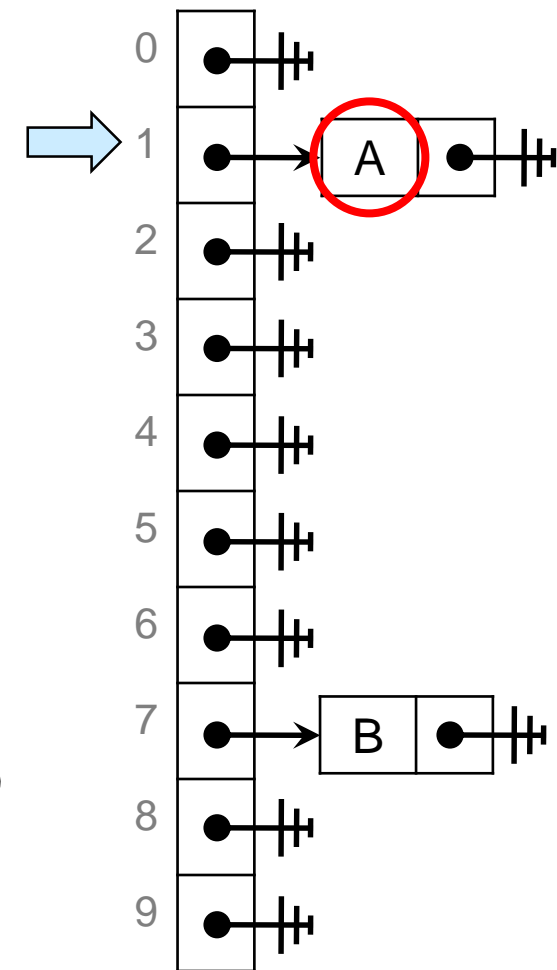
-1189657311

Ok. The hash index is 1.
It points to a node
for entry (A)

What's the key of (A)?

"apple"

Is it the same as "pumpkin"?



- ✓ new dictionary
- ✓ insert A = ("apple", 20)
- ✓ insert B = ("banana", 10)

Next, you *insert*
C = ("pumpkin", 50)

- Why is the library asking this?
 - it does not know the type of keys
 - even if it did, there are many ways to compare them
- *You need to tell it*

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

- ✓ new dictionary
- ✓ insert A = ("apple", 20)
- ✓ insert B = ("banana", 10)

Next, you *insert*
C = ("pumpkin", 50)

Client Implementation

Insert C = ("pumpkin", 50)

What's the key of (C)?

"pumpkin"

What's its hash value?

-1189657311

Ok. The hash index is 1.
It points to a node
for entry (A)

What's the key of (A)?

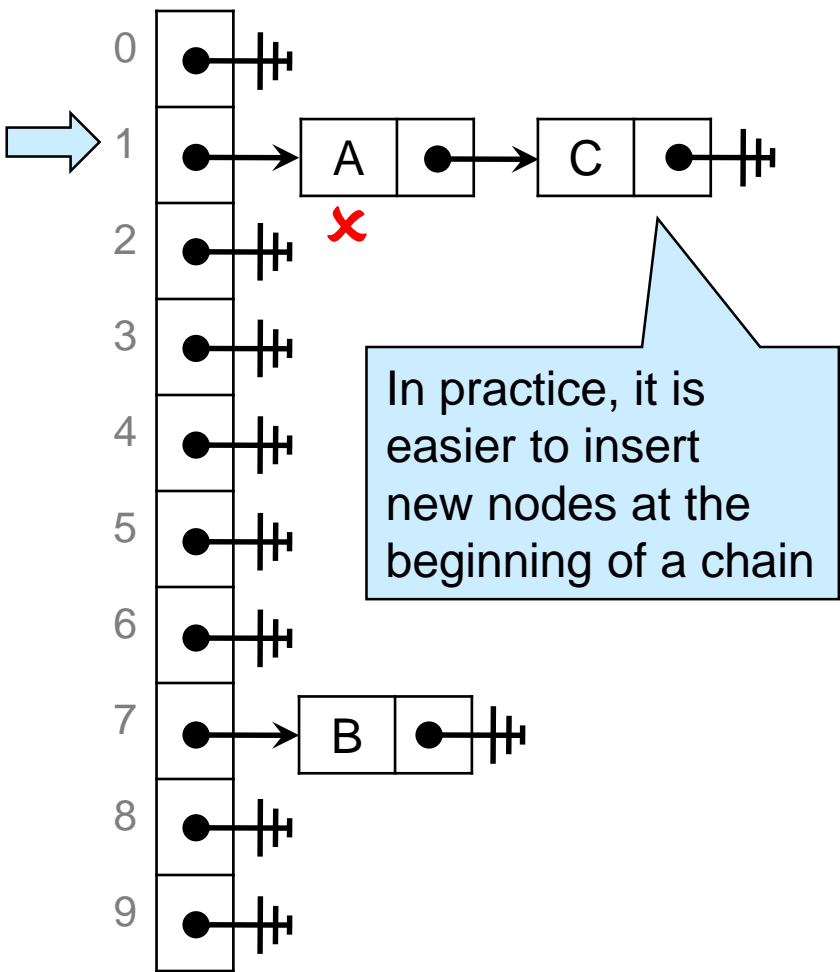
"apple"

Is it the same as "pumpkin"?

No

There is no next node.
I can insert entry
(C) there.

Done



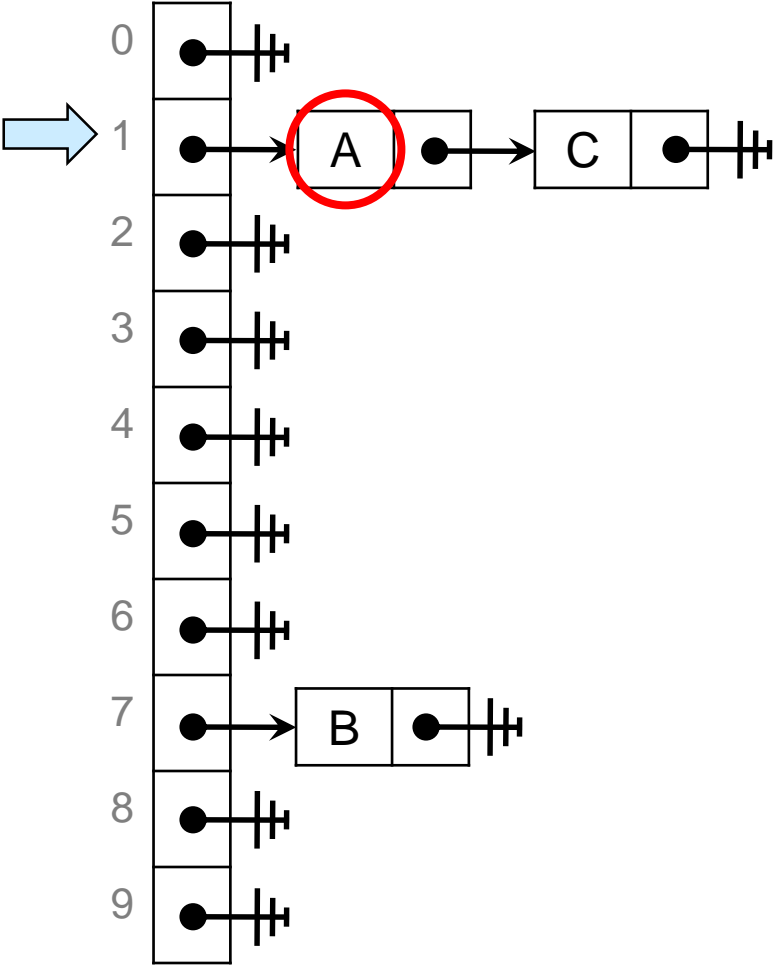
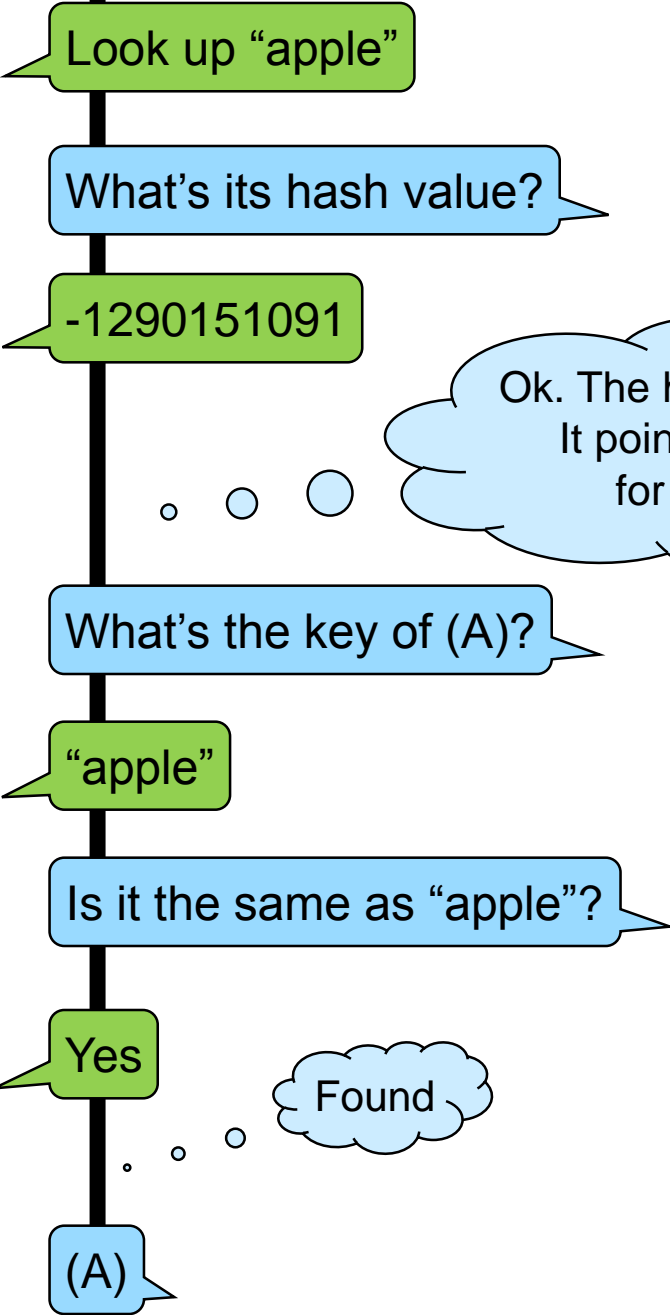
- The library asked for
 - the key of the entry
 - the hash value of the key
 - whether two keys are the same

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

- ✓ new dictionary
- ✓ insert A = ("apple", 20)
- ✓ insert B = ("banana", 10)
- ✓ insert C = ("pumpkin", 50)

Next, you
look up "apple"

Client Implementation



- Looking up a key follows the same steps as inserting an entry

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

Client Implementation

Look up "lime"

What's its hash value?

2086736531

Ok. The hash index is 1.
It points to a node
for entry (A)

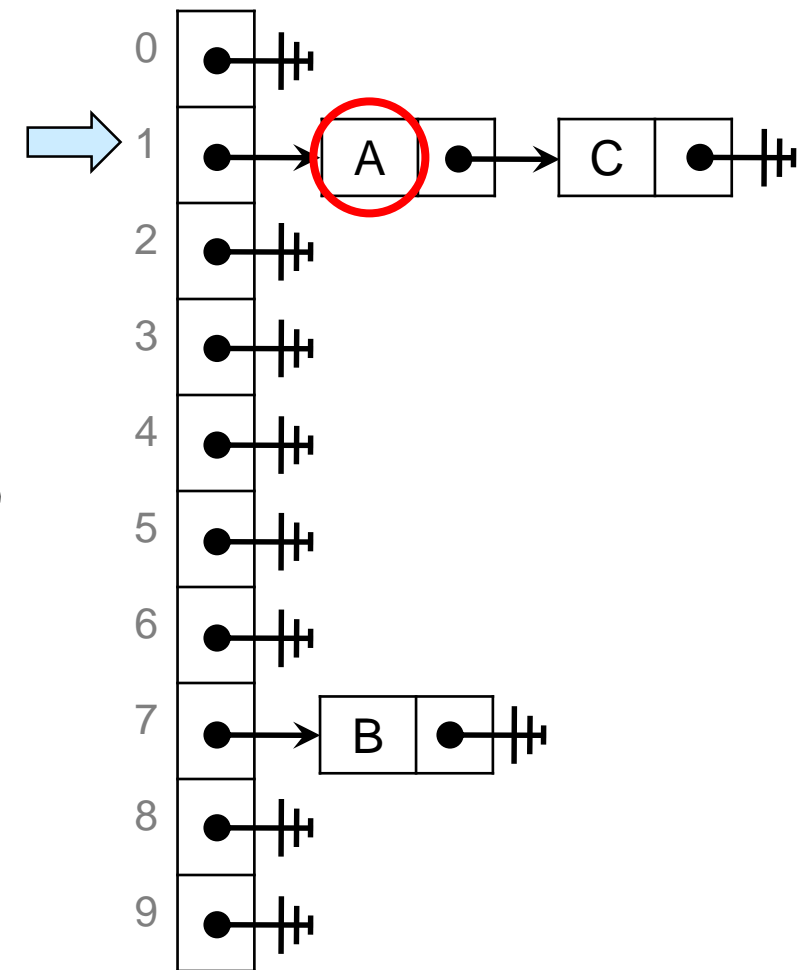
What's the key of (A)?

"apple"

Is it the same as "lime"?

- ✓ new dictionary
- ✓ insert A = ("apple", 20)
- ✓ insert B = ("banana", 10)
- ✓ insert C = ("pumpkin", 50)
- ✓ look up "apple"

*Next, you
look up "lime"*



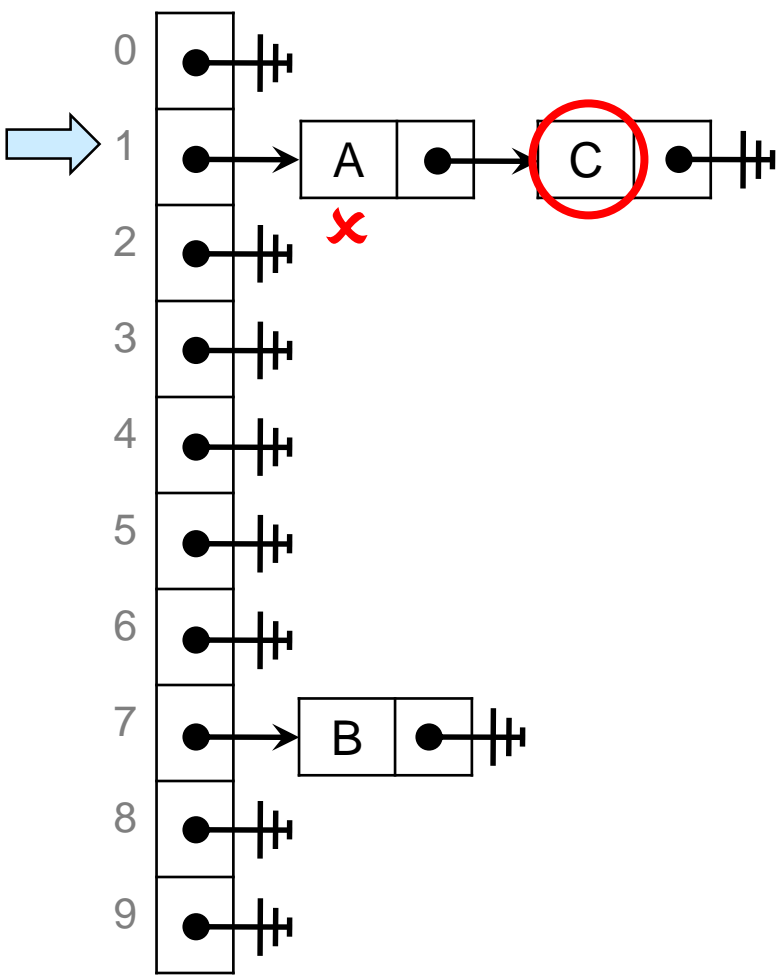
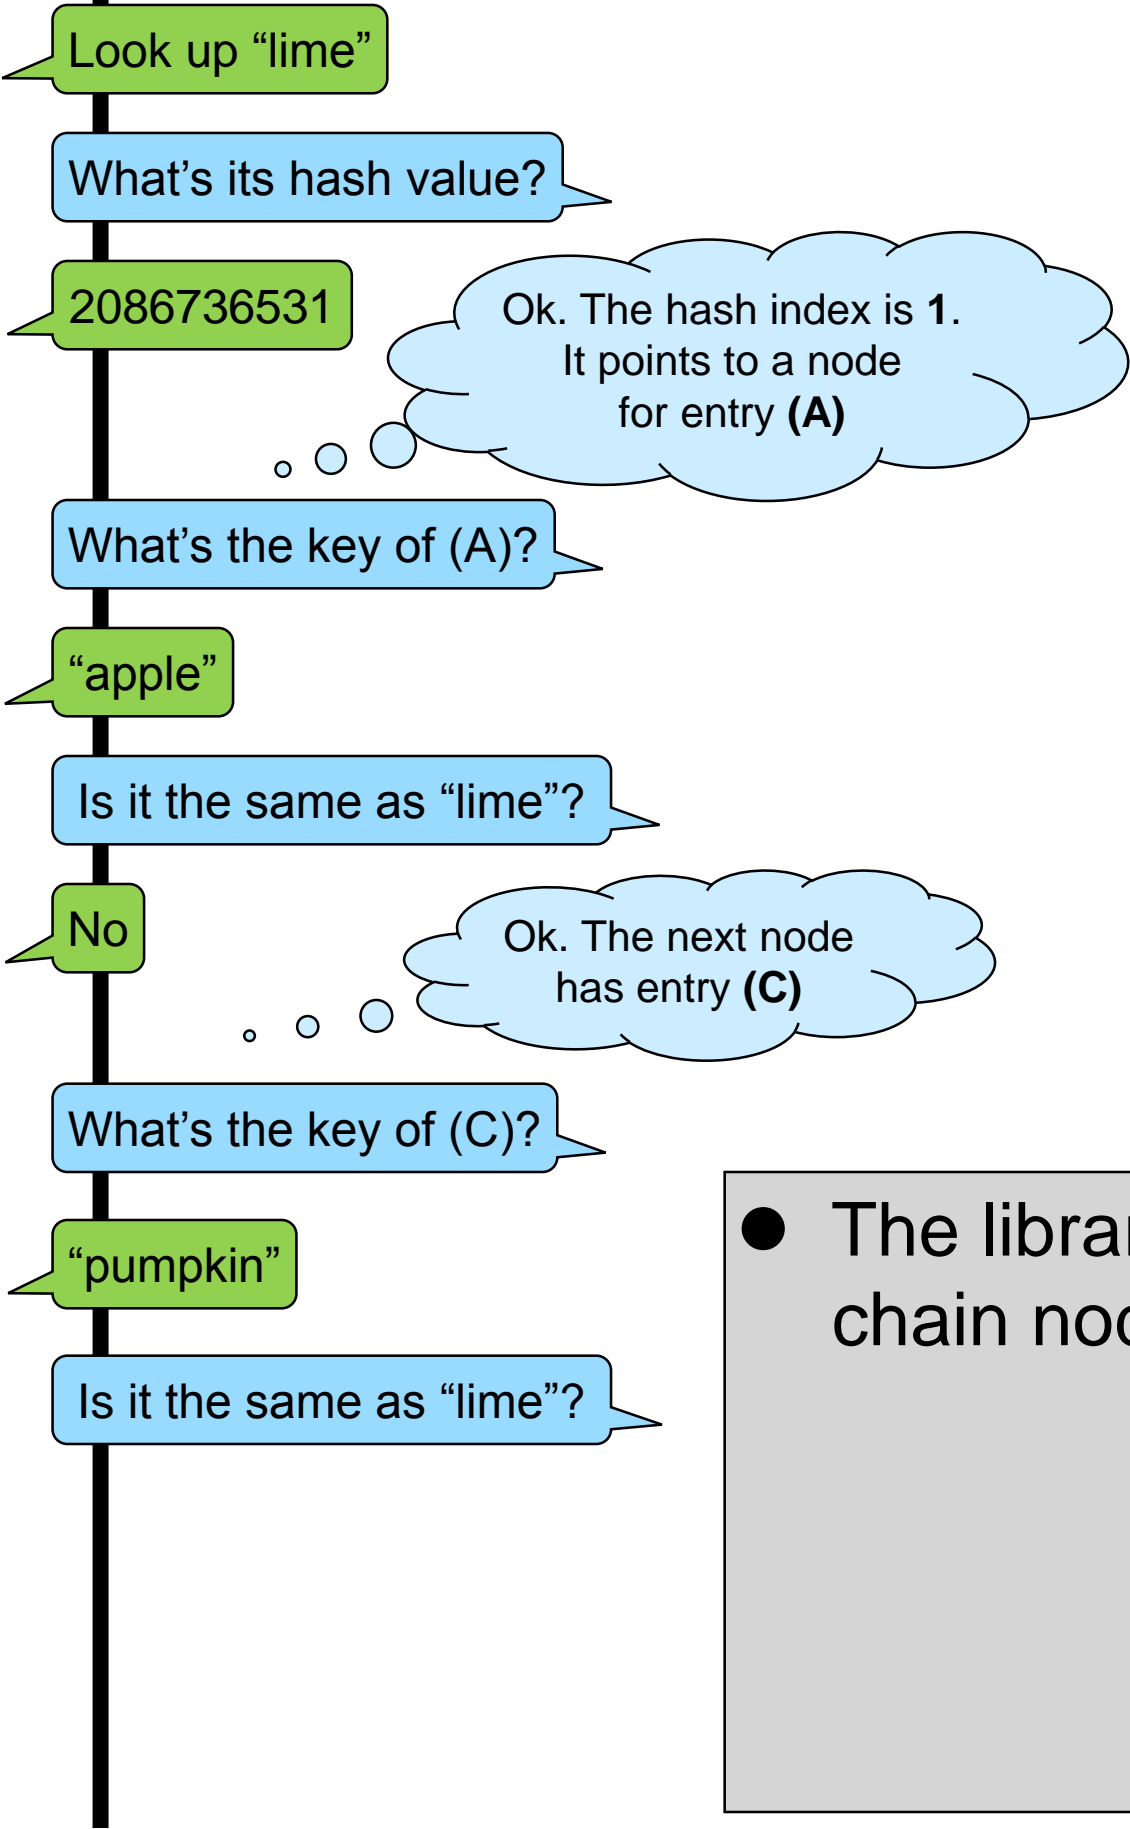
- The library goes through the chain node by node

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

- ✓ new dictionary
- ✓ insert A = ("apple", 20)
- ✓ insert B = ("banana", 10)
- ✓ insert C = ("pumpkin", 50)
- ✓ look up "apple"

Next, you
look up "lime"

Client Implementation



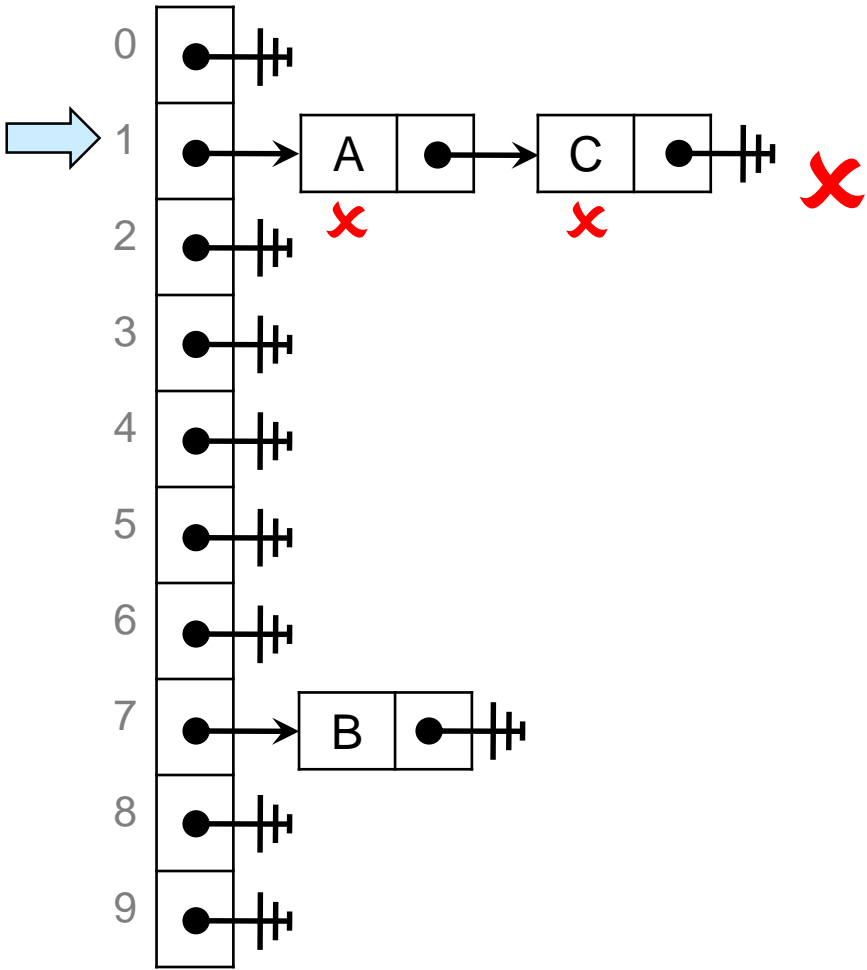
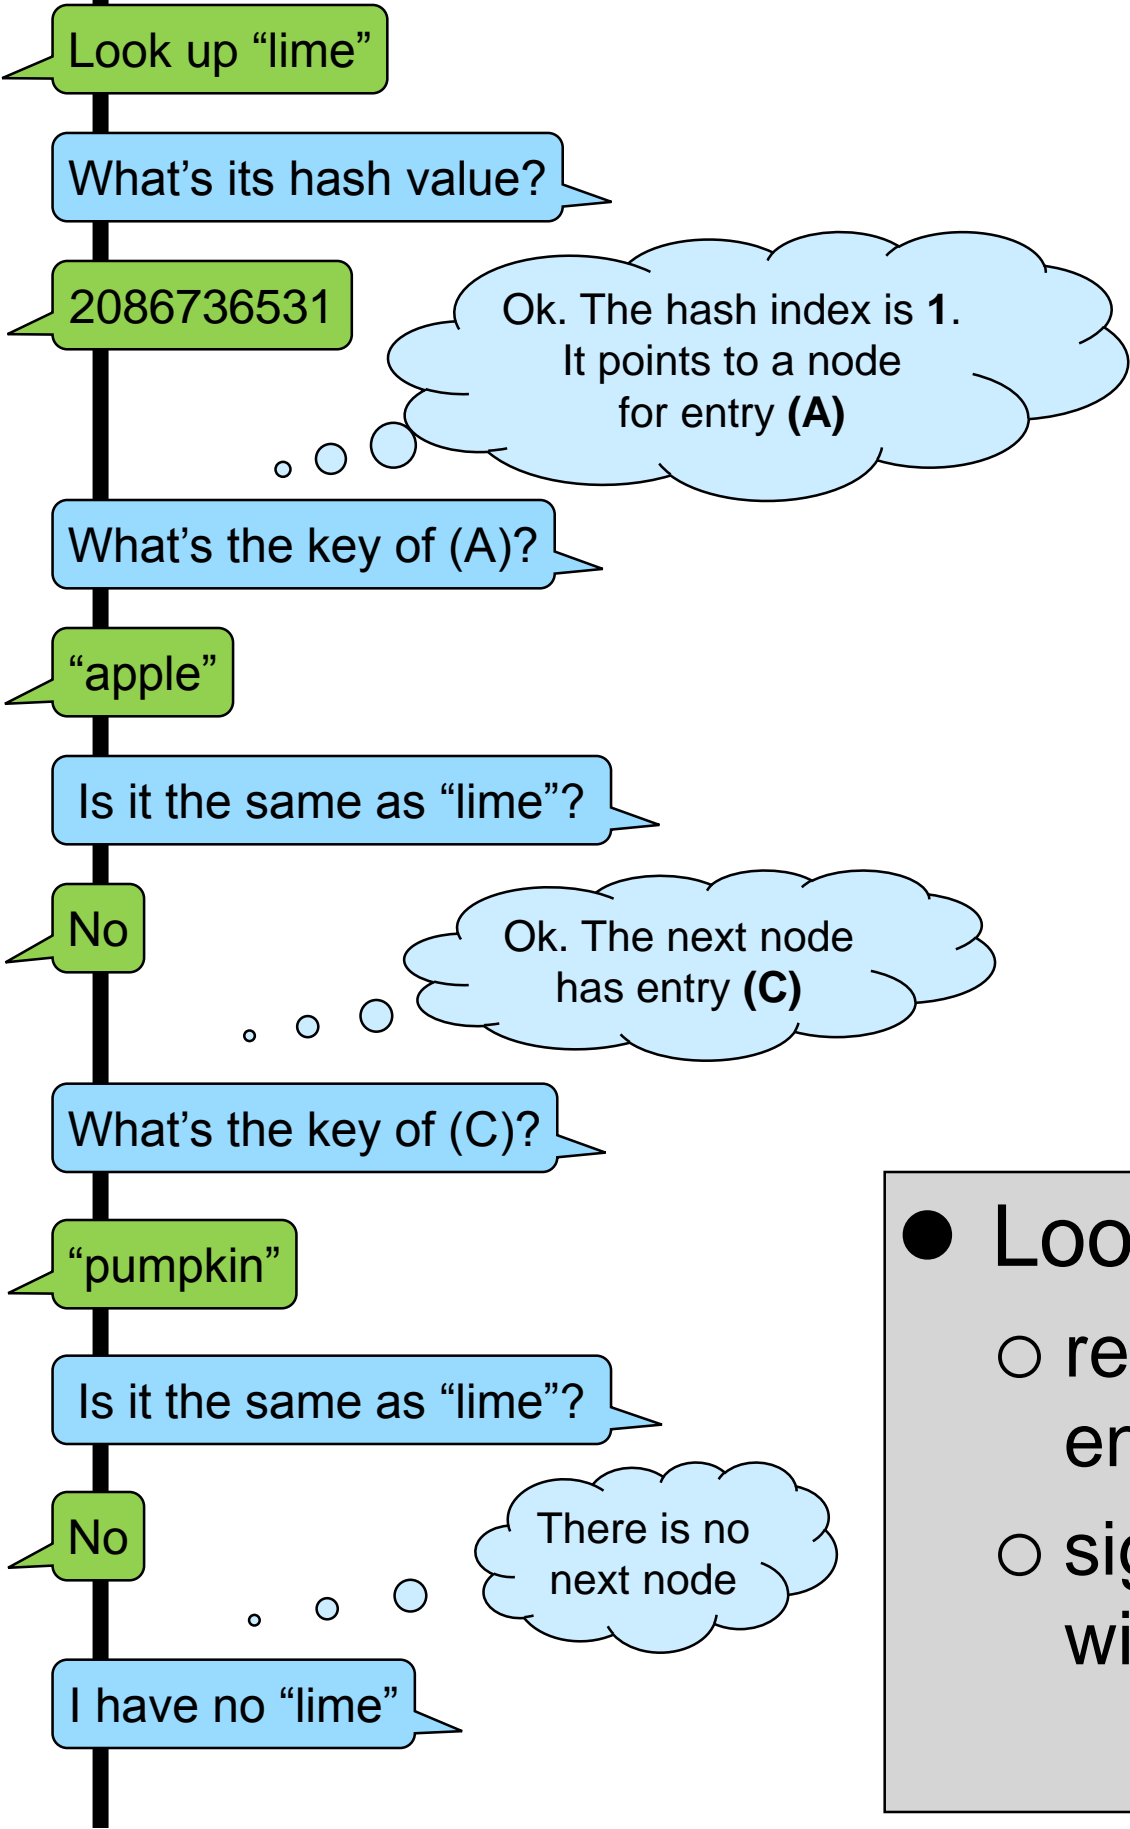
- The library goes through the chain node by node

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

- ✓ new dictionary
- ✓ insert A = ("apple", 20)
- ✓ insert B = ("banana", 10)
- ✓ insert C = ("pumpkin", 50)
- ✓ lookup "apple"

Next, you
look up "lime"

Client Implementation



- Looking up a key can
 - return the associated entry, or
 - signal there is no entry with this key

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

- ✓ new dictionary
- ✓ insert A = ("apple", 20)
- ✓ insert B = ("banana", 10)
- ✓ insert C = ("pumpkin", 50)
- ✓ look up "apple"
- ✓ look up "lime"

Next, you *insert*
D = ("banana", 20)

Client Implementation

Insert D = ("banana", 20)

What's the key of (B)?

"banana"

What's its hash value?

207055587

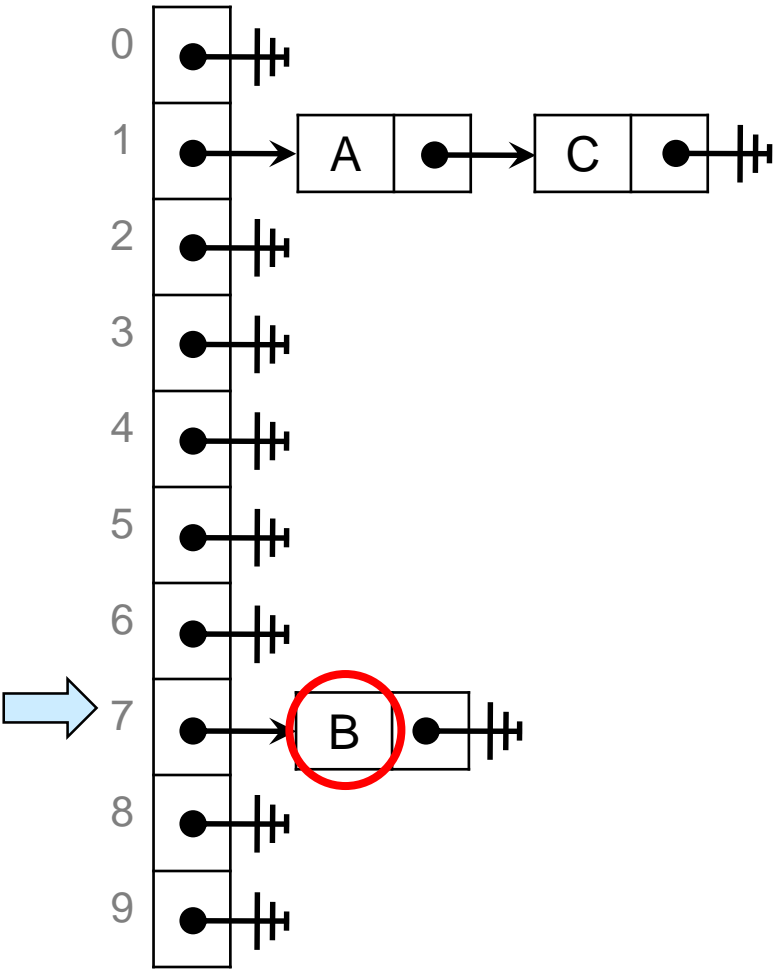
Ok. The hash index is 7.
It points to a node
for entry (B)

What's the key of (B)?

"banana"

Is it the same as "banana"?

Yes



● What to do if the key is already there?

Key	Hash
"apple"	-1290151091
"berry"	-514151789
"banana"	207055587
"grape"	-581390202
"lemon"	-665562942
"lime"	2086736531
"pumpkin"	-1189657311

- ✓ new dictionary
- ✓ insert A = ("apple", 20)
- ✓ insert B = ("banana", 10)
- ✓ insert C = ("pumpkin", 50)
- ✓ look up "apple"
- ✓ look up "lime"

Next, you *insert*
D = ("banana", 20)

Client Implementation

Insert D = ("banana", 20)

What's the key of (B)?

"banana"

What's its hash value?

207055587

Ok. The hash index is 7.
It points to a node
for entry (B)

What's the key of (B)?

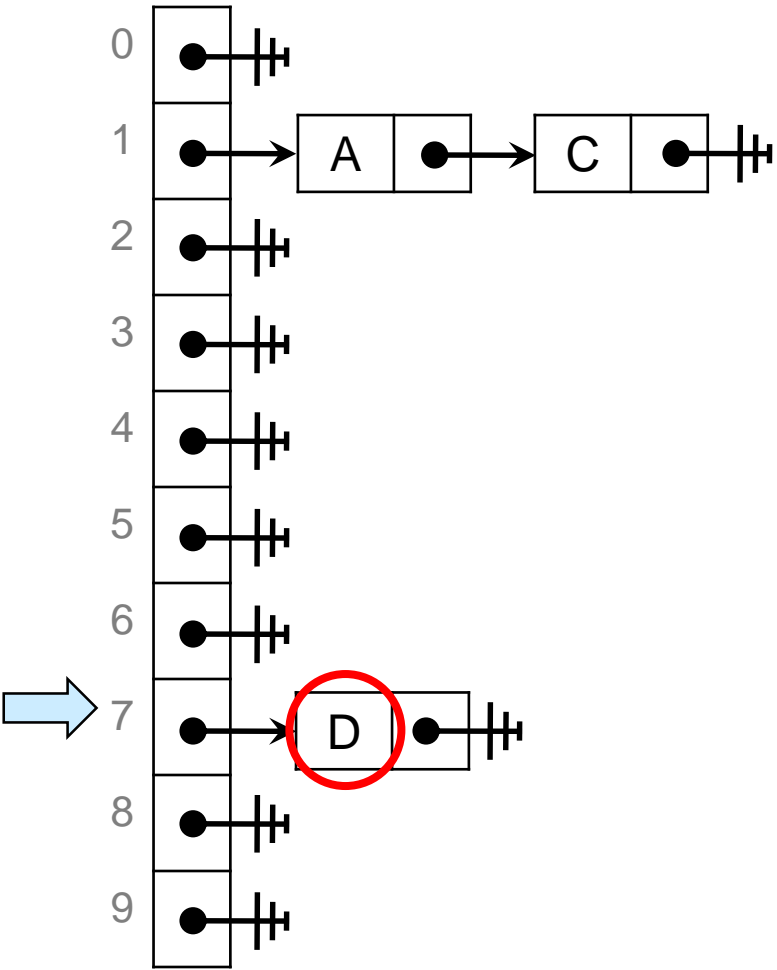
"banana"

Is it the same as "banana"?

Yes

Ok. This is where
to insert (D)

Done



- What to do if the key is already there?
- **Overwrite** the stored entry

What Have we Learned?

- The library needs information from the client to do its job
 - the key of an entry
 - the hash value of a key
 - whether two keys are the same
- How shall the client provide this information?
 - Back and forth like we did?
 - Too cumbersome
 - we want to just call lookup and get a result
 - Supply **functions** the library can use to find this information
 - a function that returns the key of an entry
 - a function that computes the hash value of a key
 - a function that determines whether two keys are the same

Hash Dictionary Interface

What the Library Provides

- A type for using dictionaries

- `hdict_t`

`hdict_t` because we will be implementing it using hash tables

- Some operations

- creating a new dictionary

- `hdict_new`

- looking up a key in a dictionary

- `hdict_lookup`

- inserting an entry into a dictionary

- `hdict_insert`

Real dictionary libraries provide many more operations.
Let's keep it simple

Let's write
the interface
of this library

Creating a Dictionary

By now we anticipate this will be a pointer ...

... and that a dictionary shall never be NULL

Library Interface

```
// typedef _____* hdict t;  
  
hdict_t hdict_new(int capacity)  
/* @requires capacity > 0; @*/  
/* @ensures \result != NULL; @*/;  
  
// ...
```

- Clients have a sense of how many entries may end up in a dictionary
 - Let them specify an initial capacity
 - whether the implementation is self-resizing or not
 - An initial capacity of 0 makes no sense
 - disallow it in precondition

Looking up a key

- `hdict_lookup` looks up a key in a dictionary ...

- we need a type `key` of keys

- ... and returns the associated entry ...

- we need a type `entry` of entries

- .. unless there is no entry with this key in the dictionary

- it then must signal that no entry was found

- Arrange so that `entry` is a **pointer type**

- either a pointer to the entry it found

- or NULL to represent “not found”

Library Interface

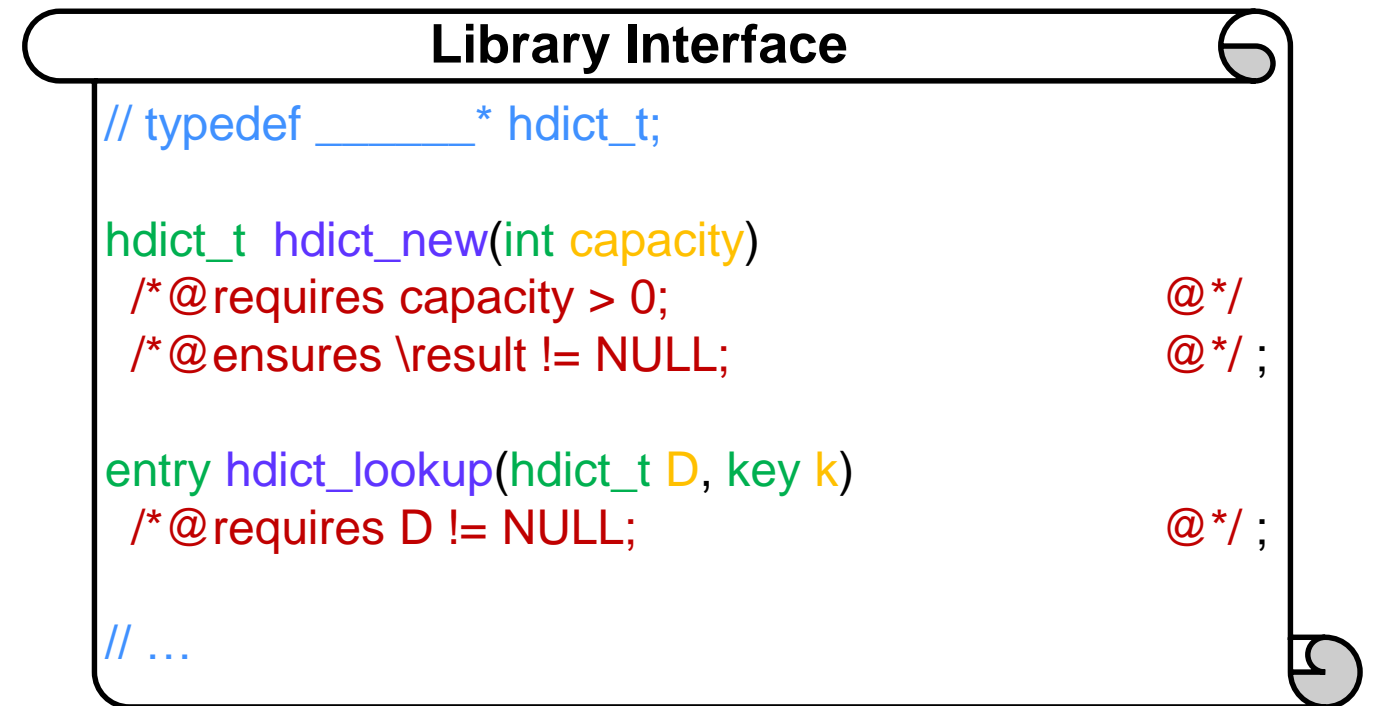
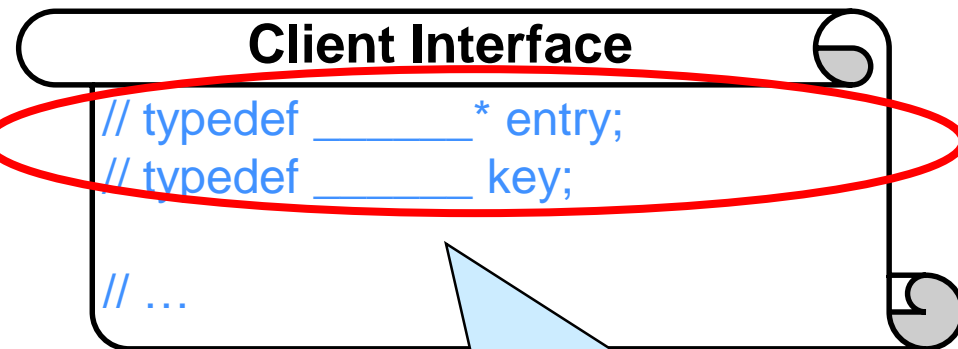
```
// typedef _____* hdict_t;

hdict_t hdict_new(int capacity)
/* @requires capacity > 0;                @*/
/* @ensures \result != NULL;              @*/;

entry hdict_lookup(hdict_t D, key k)
/* @requires D != NULL;                    @*/;

// ...
```

Key and Entry Types



- It's the client who decides what keys and entries are
 - the interface must tell the client to do this
- The interface has two parts
 - the **client interface**: what the client needs to supply to the library
 - the **library interface**: what the library provides to the client

Inserting an Entry

Client Interface

```
// typedef _____* entry;  
// typedef _____ key;  
  
// ...
```

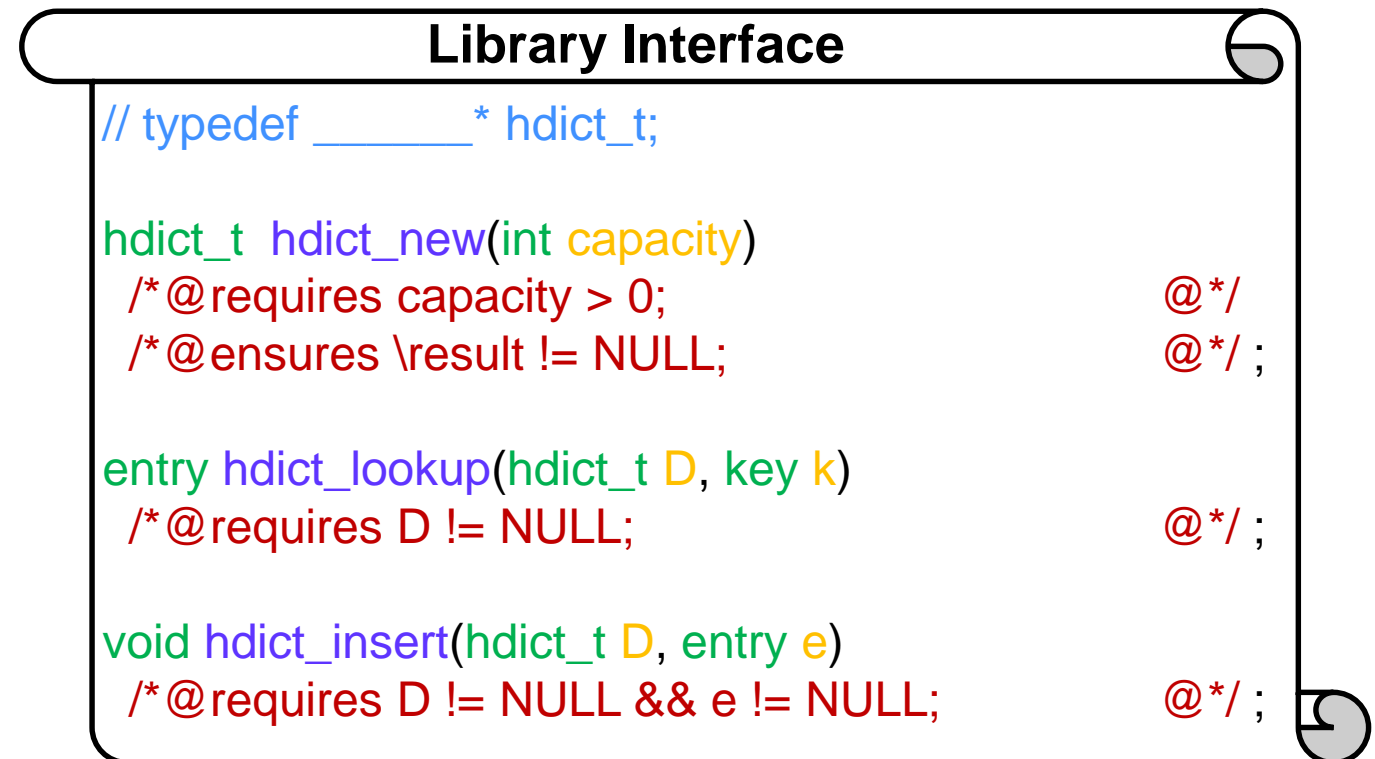
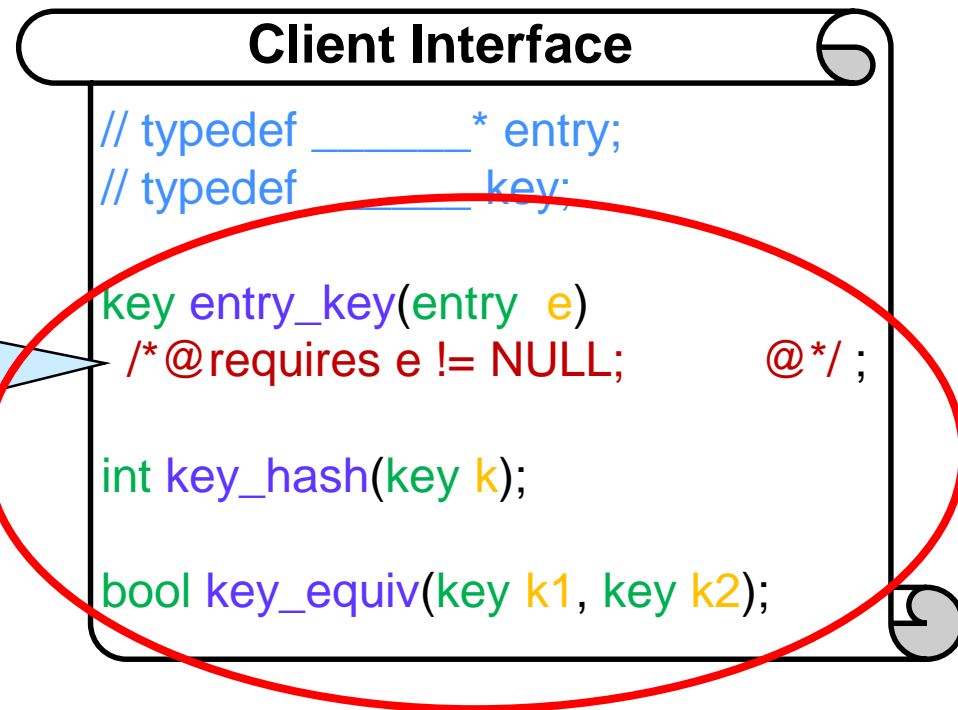
Library Interface

```
// typedef _____* hdict_t;  
  
hdict_t hdict_new(int capacity)  
/* @requires capacity > 0; @*/  
/* @ensures \result != NULL; @*/;  
  
entry hdict_lookup(hdict_t D, key k)  
/* @requires D != NULL; @*/;  
  
void hdict_insert(hdict_t D, entry e)  
/* @requires D != NULL && e != NULL; @*/;
```

- If NULL stands for an entry that was not found, no entry shall ever be NULL

e cannot be NULL

What about all those Questions?



- The library needs information from the client to do its job
 - Supply **functions** the library can use to find this information
 - a function that returns the key of an entry entry_key
 - a function that computes the hash value of a key key_hash
 - a function that determines whether two keys are the same key_equiv
 - Add their prototype in the client interface!

A Postcondition for `hdict_insert`

Client Interface

```
// typedef _____* entry;
// typedef _____ key;

key entry_key(entry e)
/*@requires e != NULL;    @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);
```

Library Interface

```
// typedef _____* hdict_t;

hdict_t hdict_new(int capacity)
/*@requires capacity > 0;    @*/
/*@ensures \result != NULL;  @*/;

entry hdict_lookup(hdict_t D, key k)
/*@requires D != NULL;    @*/;

void hdict_insert(hdict_t D, entry e)
/*@requires D != NULL && e != NULL;    @*/
/*@ensures hdict_lookup(D, entry_key(e)) == e; @*/;
```

- If we insert an entry and lookup its key, we should find that entry
 - i.e., `hdict_lookup(D, entry_key(e)) == e`
 - lookup returns the very entry `e`
 - not a different entry with the same data

e is a pointer

A Postcondition for `hdict_lookup`

```
Client Interface

// typedef _____* entry;
// typedef _____ key;

key entry_key(entry e)
/* @requires e != NULL;    @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);
```

```
Library Interface

// typedef _____* hdict_t;

hdict_t hdict_new(int capacity)
/* @requires capacity > 0;    @*/
/* @ensures \result != NULL;  @*/;

entry hdict_lookup(hdict_t D, key k)
/* @requires D != NULL;    @*/
/* @ensures \result == NULL
|| key_equiv(entry_key(\result), k);    @*/;

void hdict_insert(hdict_t D, entry e)
/* @requires D != NULL && e != NULL;    @*/
/* @ensures hdict_lookup(D, entry_key(e)) == e; @*/;
```

- If we look up a key
 - either we get back NULL
 - `\result == NULL`
 - or the key of the returned entry is our key
 - `key_equiv(entry_key(\result), k)`
- The client interface functions give us a way to write very precise postconditions

The Hash Dictionary Interface

Client Interface

```
// typedef _____* entry;
// typedef _____ key;

key entry_key(entry e)
/* @requires e != NULL;          @*/ ;

int key_hash(key k);

bool key_equiv(key k1, key k2);
```

What the library needs
from the client

Library Interface

```
// typedef _____* hdict_t;

hdict_t hdict_new(int capacity)
/* @requires capacity > 0;          @*/
/* @ensures \result != NULL;        @*/ ;

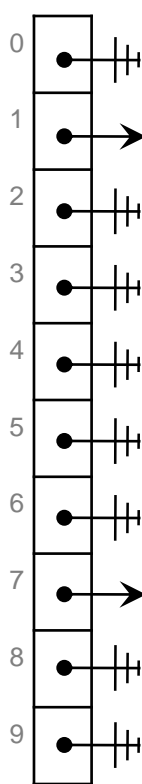
entry hdict_lookup(hdict_t D, key k)
/* @requires D != NULL;              @*/
/* @ensures \result == NULL
           || key_equiv(entry_key(\result), k); @*/ ;

void hdict_insert(hdict_t D, entry e)
/* @requires D != NULL && e != NULL; @*/
/* @ensures hdict_lookup(D, entry_key(e)) == e; @*/ ;
```

What the library provides
to the client

Hash Dictionary Implementation

Hash Dictionary Types



```
typedef struct chain_node chain;
struct chain_node {
    entry data;    // data != NULL
    chain* next;
};

struct hdict_header {
    int size;      // size >= 0
    int capacity;  // capacity > 0
    chain*[] table; // \length(table) == capacity
};
typedef struct hdict_header hdict;

// ... rest of implementation

typedef hdict* hdict_t;
```

Implementation

These are expected constraints on the fields

As usual, the abstract client type is a pointer to the concrete implementation type

- Each chain is a NULL-terminated linked list of entries
 - entries can't be NULL
- A dictionary is implemented as a hash table
- We need to keep track of
 - size: the number of entries
 - capacity: the length of the hash table
 - the hash table itself
 - an array of pointers to chain nodes

Representation Invariants

```
typedef struct hdict_header hdict;  
struct hdict_header {  
    int size;           // size >= 0  
    int capacity;       // capacity > 0  
    chain*[] table;     // \length(table) == capacity  
};
```

- We need to capture the field constraints in the type

```
bool is_array_expected_length(chain*[] A, int len) {  
    //@assert \length(A) == len;  
    return true;  
}  
  
// Representation invariant  
bool is_hdict(hdict* H) {  
    return H != NULL  
        && H->size >= 0  
        && H->capacity > 0  
        && is_array_expected_length(H->table, H->capacity);  
}  
  
// ... rest of implementation
```

Implementation

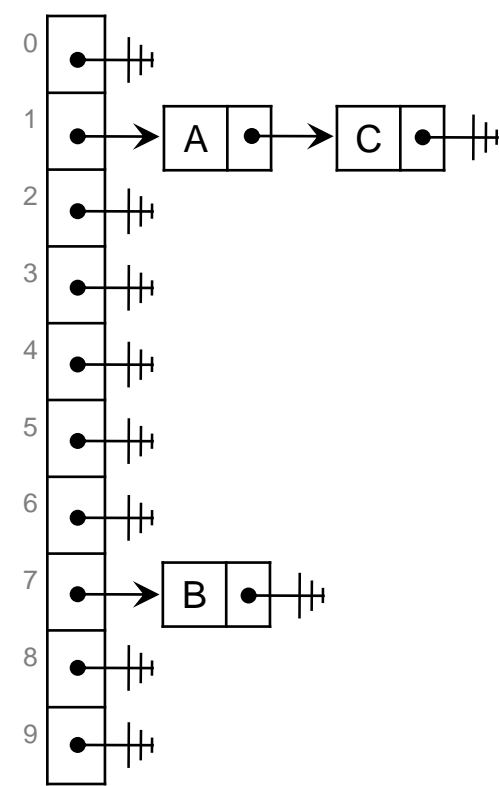
Usual trick to check the length of an array

Abstract data structures are never NULL

Field constraints

More Representation Invariants

- Hash tables have a much more involved structure than previous concrete library types
 - the chains are acyclic
 - no two entries have the same key
 - each entry hashes to the right index
 - no entry is NULL
 - the number of entries equals the size field



```
// Representation invariant
bool is_hdict(hdict* H) {
    return H != NULL
        && H->size >= 0
        && H->capacity > 0
        && is_array_expected_length(H->table, H->capacity)
        && is_valid_hashtable(H);
}
```

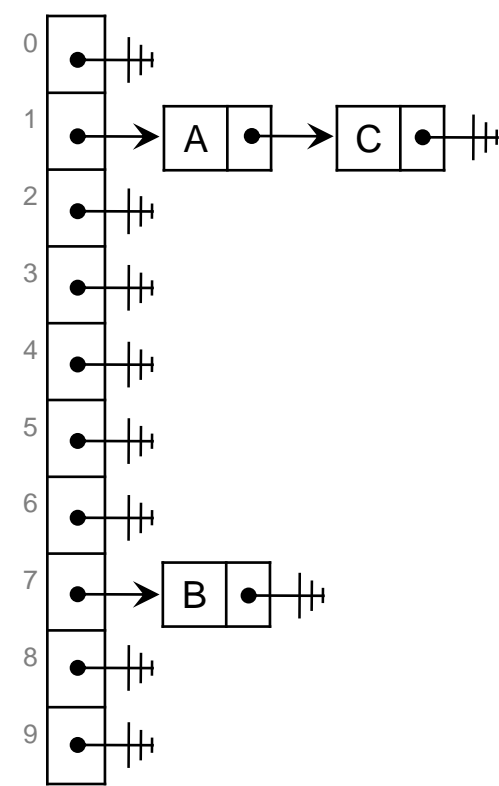
Implementation

This tests all these structural constraints

Exercise!

Invalidating Invariants

- The client can modify the keys **after** they have been inserted in the hash table
 - The chains contain **pointers** to entries
- This can invalidate the data structure invariants
 - `is_dict` fails the next time it is called
 - this is not because of a bug in the library
 - this is because the client manipulated the entries through aliases



- Aliasing is dangerous!

This couldn't happen in
any of the data structures
we studied so far

Implementing `hdict_lookup`

Library Interface

```
entry hdict_lookup(hdict_t D, key k)
/* @requires D != NULL;                                     @*/
/* @ensures \result == NULL
    || key_equiv(entry_key(\result), k); @*/;
```

Client Interface

```
// typedef _____* entry;
// typedef _____ key;

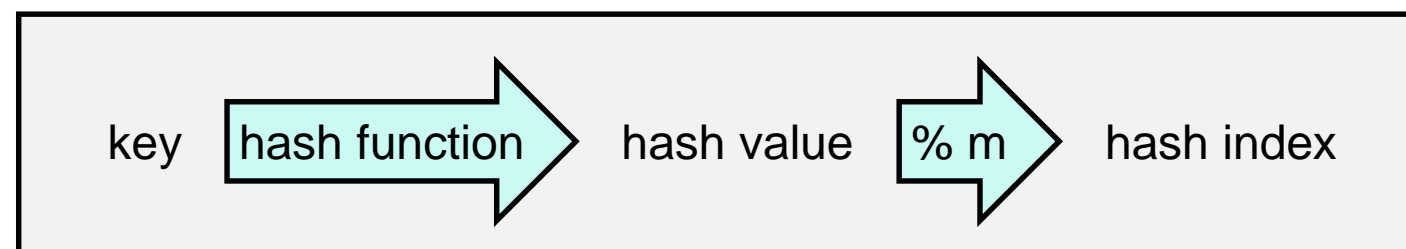
key entry_key(entry e)
/* @requires e != NULL; @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);
```

- First we need to find the right bucket

- determine the hash index of `k`



```
int i = key_hash(k) % D->capacity;
```

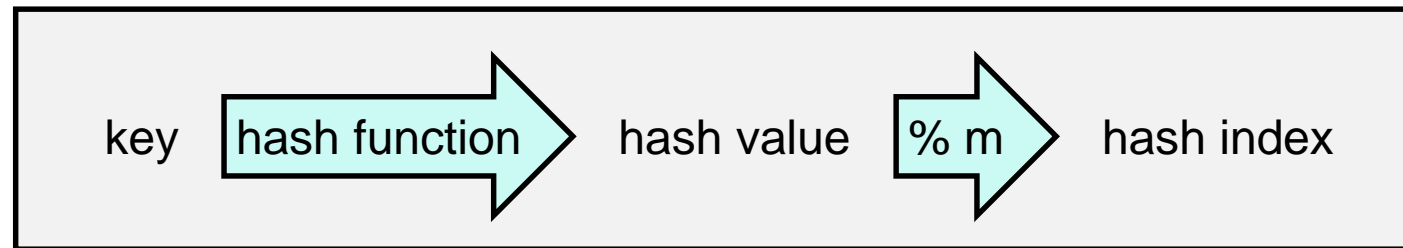
What's its hash value?

The library's questions are answered by the client interface functions

- This won't work if `hash_key(k)` is negative!

```
int i = abs(key_hash(k) % D->capacity);
```

Finding the Right Bucket



Client Interface

```
// typedef _____* entry;
// typedef _____ key;

key entry_key(entry e)
/*@requires e != NULL; @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);
```

➤ determine the hash index of **k**

```
int i = abs(key_hash(k) % H->capacity);
```

- We will need to do the same in **hdict_insert**
 - **factor it out** in a function that computes the hash index of a key

```
int index_of_key(hdict* H, key k)
/*@requires is_hdict(H);
/*@ensures 0 <= \result && \result < H->capacity;
{
    return abs(key_hash(k) % H->capacity);
}
```

Implementation

Implementing `hdict_lookup`

- First we need to find the right bucket
- Then we go through its chain
 - extract the key of each entry
 - check if it is equal to `k`

Client Interface

```
// typedef _____* entry;
// typedef _____ key;

key entry_key(entry e)
/*@requires e != NULL;  @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);
```

Library Interface

```
entry hdict_lookup(hdict_t D, key k)
/*@requires D != NULL;  @*/
/*@ensures \result == NULL
|| key_equiv(entry_key(\result), k); @*/;
```

What's its key?

Is it the same as `k`?

```
entry hdict_lookup(hdict* H, key k)
/*@requires is_hdict(H);
/*@ensures \result == NULL
|| key_equiv(entry_key(\result), k);
{
  int i = index_of_key(H, k);
  for (chain* p = H->table[i]; p != NULL; p = p->next) {
    if (key_equiv(entry_key(p->data), k))
      return p->data;
  }
  return NULL;
}
```

Implementation

H must satisfy the representation invariant

`i` is the hash index of `k`

This is the start of the chain

Return the entry if `k` is found ...

... and signal it's not there otherwise

Implementing `hdict_insert`

Client Interface

```
// typedef _____* entry;
// typedef _____ key;

key entry_key(entry e)
/*@requires e != NULL;  @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);
```

Library Interface

```
void hdict_insert(hdict_t D, entry e)
/*@requires D != NULL && e != NULL;  @*/
/*@ensures hdict_lookup(D, entry_key(e)) == e;  @*/;
```

```
void hdict_insert(hdict* H, entry e)
/*@requires is_hdict(H) && e != NULL;
/*@ensures \ hdict_lookup(D, entry_key(e)) == e;
/*@ensures is_hdict(H);
```

```
{
    key k = entry_key(e);
    int i = index_of_key(H, k);
    for (chain* p = H->table[i]; p != NULL; p = p->next) {
        if (key_equiv(entry_key(p->data), k)) {
            p->data = e;
            return;
        }
    }
    chain* p = alloc(chain);
    p->data = e;
    p->next = H->table[i];
    H->table[i] = p;
    (H->size)++;
}
```

Implementation

H must remain valid after the insertion

Check if there is already an entry with the same key (similar code to `hdict_lookup`)

If so overwrite it

Otherwise, prepend a new node containing `e`

Implementing `hdict_new`

```
hdict* hdict_new(int capacity)
/*@requires capacity > 0;
/*@ensures is_hdict(\result);
{
    hdict* H = alloc(hdict);
    H->size = 0;
    H->capacity = capacity;
    H->table = alloc_array(chain*, capacity);
    return H;
}
```

Implementation

Returned dictionary must be valid

Initialized to default pointer value, NULL

Client Interface

```
// typedef _____* entry;
// typedef _____ key;

key entry_key(entry e)
/*@requires e != NULL;  @*/;

int key_hash(key k);

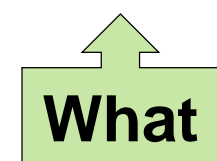
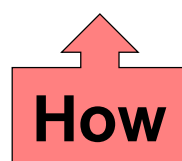
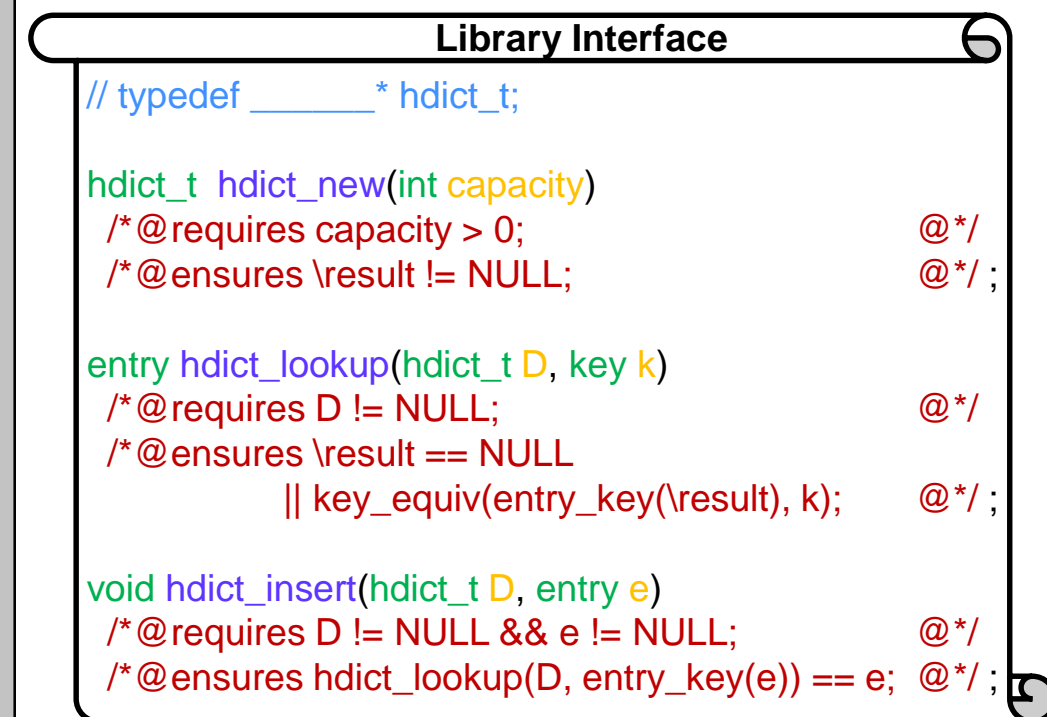
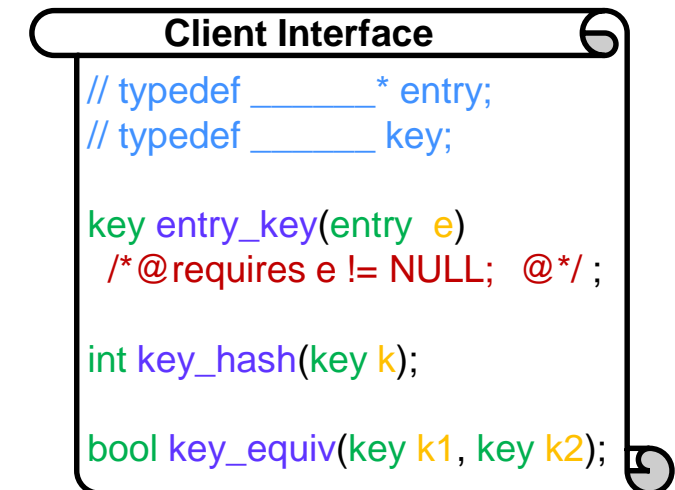
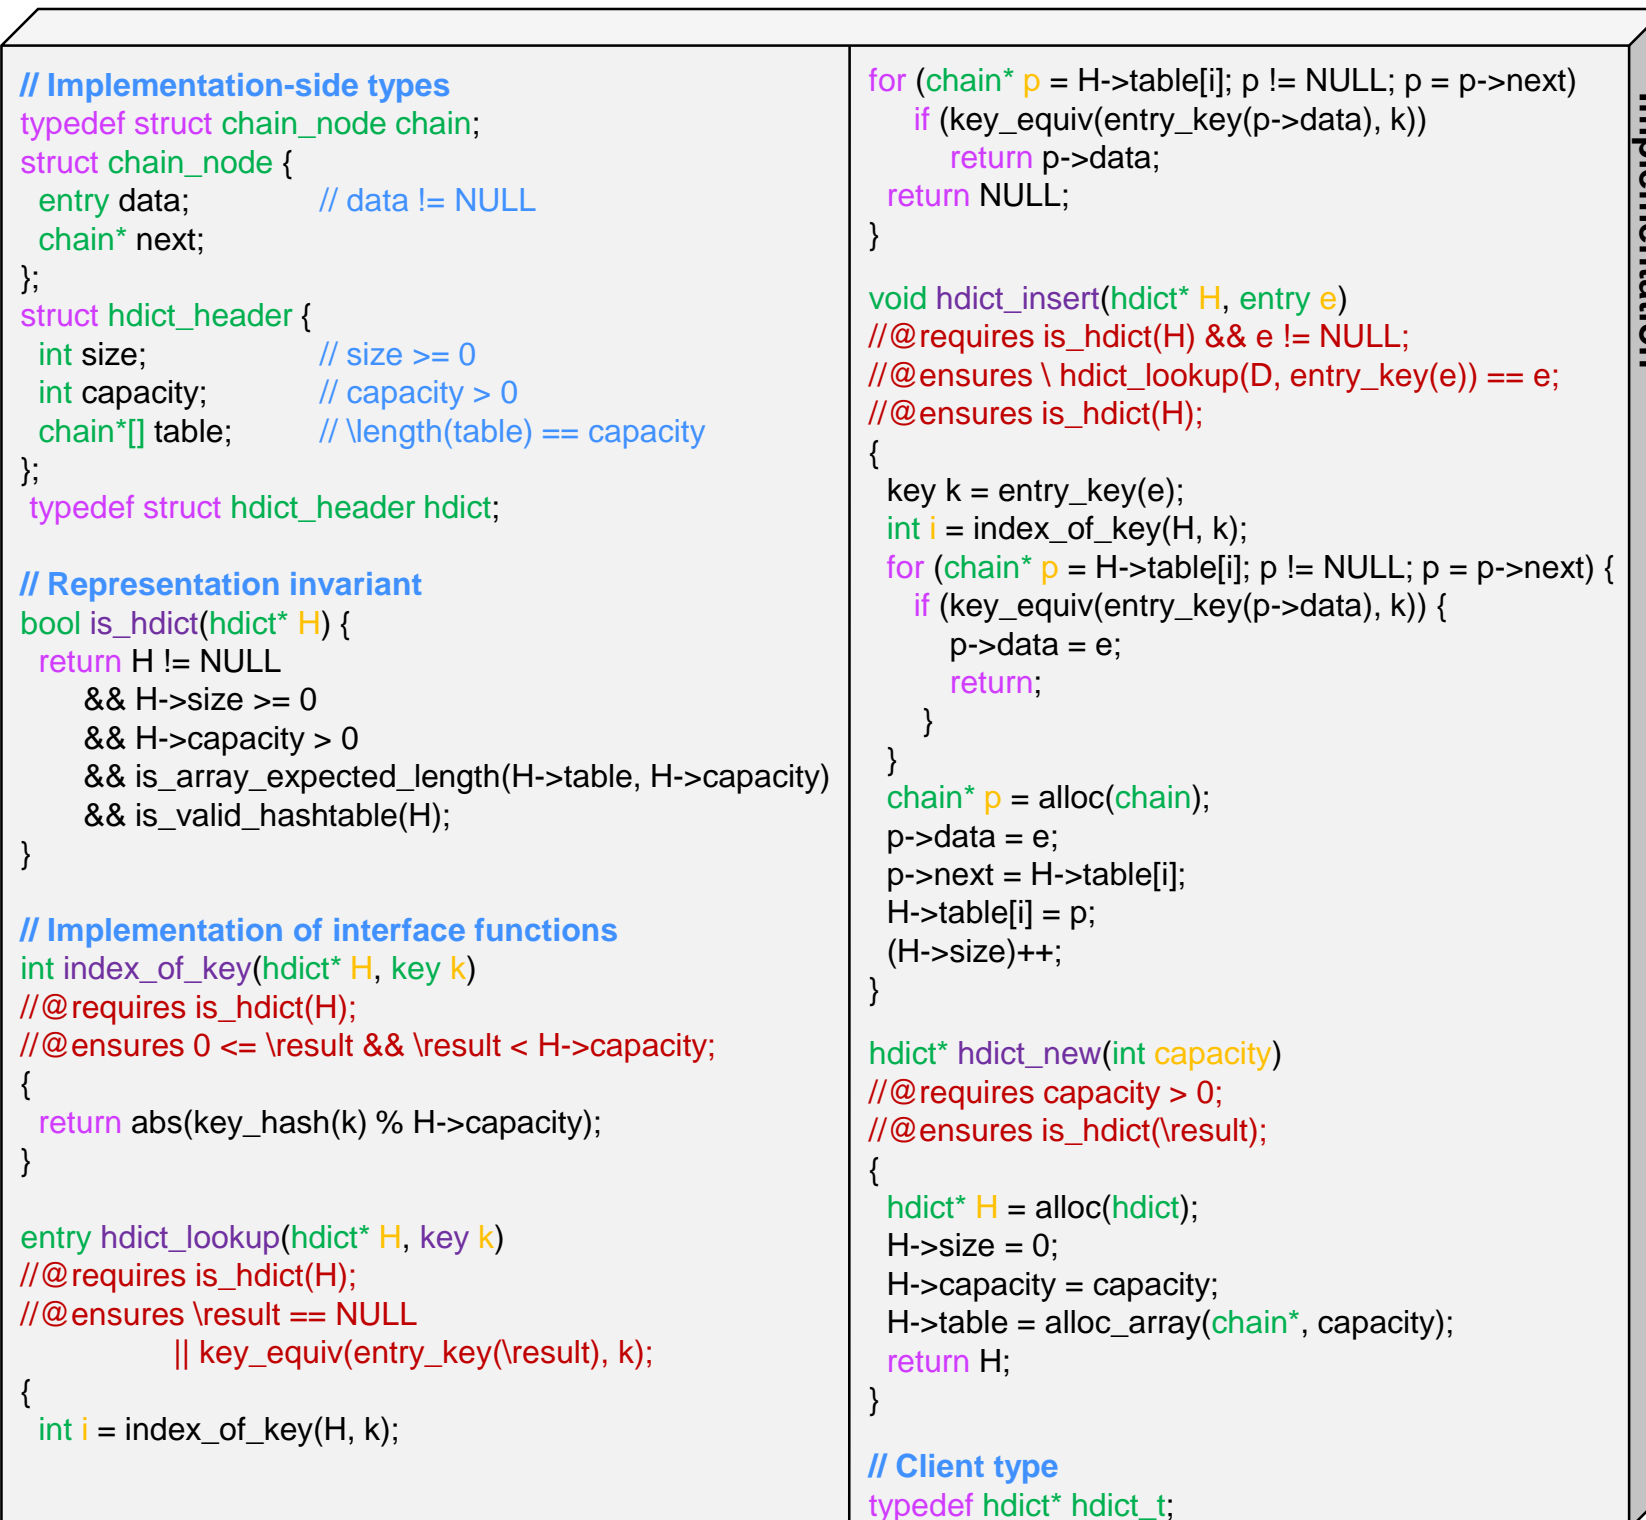
bool key_equiv(key k1, key k2);
```

Library Interface

```
hdict_t hdict_new(int capacity)
/*@requires capacity > 0;  @*/
/*@ensures \result != NULL; @*/;
```

The Hash Dictionary Library

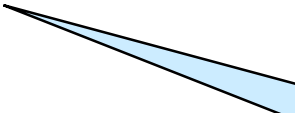
Overall Implementation



Complex Libraries

- The hash dictionary library is a **complex library**


- it needs the client to supply code and functions
- so that it can provide its services



Stacks and queues were **elementary libraries**

- Complex libraries consist of


- a **client interface**
- an **implementation**
- a **library interface**



They consisted of only an implementation and a library interface

- The client sees the client and library interfaces

- but not the implementation



Their client only saw the library interface

Structure of a Complex C0 Library File

NEW

```
/****** CLIENT INTERFACE *****/  
// typedef _____ *entry;  
...  
key entry_key(entry e)  
/*@requires e != NULL;          @*/;  
...
```

Client
interface

- Client interface
 - Client type names
 - Prototype of client functions

```
/****** IMPLEMENTATION *****/  
// Implementation-side types  
struct hdict_header {...};  
typedef struct hdict_header hdict;  
  
// Representation invariant  
bool is_hdict(hdict* H) { ... }  
  
// Implementation of interface functions  
entry hdict_lookup(hdict* H, key k)  
/*@requires is_hdict(H);          @*/  
/*@ensures ....                  @*/  
{ ... }  
...  
// Client type  
typedef hdict* hdict_t;  
  
/****** LIBRARY INTERFACE *****/  
// typedef _____ *hdict_t;  
  
entry hdict_lookup(hdict_t D, key k)  
/*@requires D != NULL;          @*/  
/*@ensures ....                  @*/;  
...
```

Implementation
Library
interface

- Implementation
 - Concrete type definition
 - Representation invariant function
 - Implementation of interface functions
 - Actual abstract type definition
- Library interface
 - Abstract type name
 - Prototype of exported functions

Structure of a Complex C0 Library File

```
/****** CLIENT INTERFACE *****/
// typedef _____ *entry;
...
key entry_key(entry e)
/*@requires e != NULL;          @*/;
...
/****** IMPLEMENTATION *****/
// Implementation-side types
struct hdict_header {...};
typedef struct hdict_header hdict;

// Representation invariant
bool is_hdict(hdict* H) { ... }

// Implementation of interface functions
entry hdict_lookup(hdict* H, key k)
/*@requires is_hdict(H);        @*/
/*@ensures ....                 @*/
{ ... }
...
// Client type
typedef hdict* hdict_t;

/****** LIBRARY INTERFACE *****/
// typedef _____ *hdict_t;

entry hdict_lookup(hdict_t D, key k)
/*@requires D != NULL;          @*/
/*@ensures ....                 @*/;
...
```

Client
interface

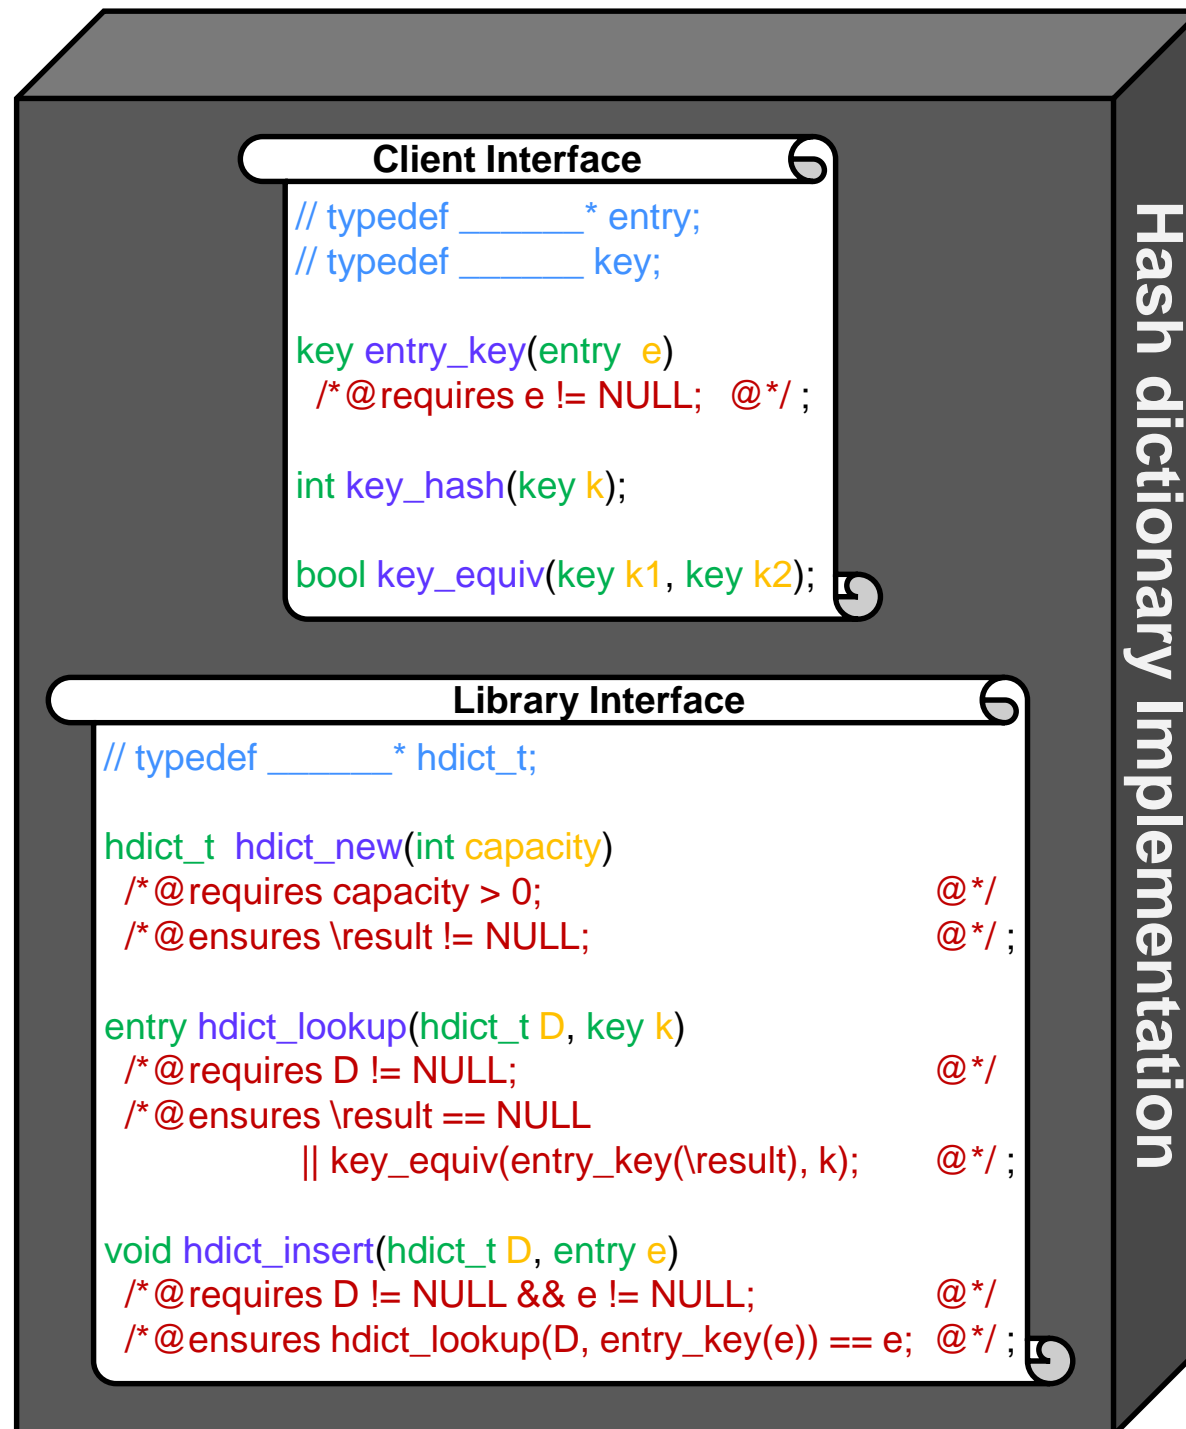
Implementation

Library
interface

- By convention,
 - the client interface is on top
 - because the implementation uses the types and functions it mentions
 - the implementation is in the middle
 - it relies on the concrete client definitions
 - it ends with the definition of the abstract client type
 - the library interface is at the bottom
 - it only mentions the abstract types

Using the Library

Using the Hash Dictionary Library



- The client needs to **define** the types and functions listed in **the client interface**
- It can **use** the types and functions exported by **the library implementation**
- The client must not rely on the implementation details

Implementing our Example

*You are the new produce manager of the local grocery store.
You want to use a dictionary to track your fruit inventory.*

- Defining the **types** requested in the client interface
 - **entries** are inventory items consisting of a fruit and a quantity
 - the fruit name is the **key**

```
// What the client wants to store in the dictionary
struct inventory_item {
    string fruit;    // key
    int quantity;
};
```

```
/****** Fulfilling the library interface *****/
typedef struct inventory_item* entry;
typedef string key;
```

This is the concrete definition of
// typedef _____* entry;

This is the concrete definition of
// typedef _____ key;

Implementing our Example

*You are the new produce manager of the local grocery store.
You want to use a dictionary to track your fruit inventory.*

- Defining the **functions** requested in client interface

```
/****** Fulfilling the library interface *****/  
key entry_key(entry e)  
//@requires e != NULL;  
{  
    return e->fruit;  
}  
  
bool key_equiv(key k1, key k2) {  
    return string_equal(k1, k2);  
}  
  
int key_hash(key k) {  
    return lcg_hash_string(k);  
}
```

The key is the fruit field
of an inventory item

Two fruit are the same
if they have the same name
(`string_equals` is defined in `<string>`)

`lcg_hash_string` is a good
hash function on strings

Client Interface Implementation

Here's the full definition of `lcg_hash_string`

- This defines every type and function in the client interface

```
Client Interface
// typedef _____* entry;
// typedef _____ key;

key entry_key(entry e)
/*@requires e != NULL; @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);
```

- We store this code in a file called **produce.c0**

Client definition file

```
#use <string>

int lcg_hash_string(string s) {
    int len = string_length(s);
    int h = 0;
    for (int i = 0; i < len; i++) {
        h = h + char_ord(string_charat(s, i));
        h = 1664525 * h + 1013904223;
    }
    return h;
}

// What the client wants to store in the dictionary
struct inventory_item {
    string fruit;      // key
    int quantity;
};

/***** Fulfilling the library interface *****/
typedef struct inventory_item* entry;
typedef string key;

key entry_key(entry e)
/*@requires e != NULL;
{
    return e->fruit;
}

bool key_equiv(key k1, key k2) {
    return string_equal(k1, k2);
}

int key_hash(key k) {
    return lcg_hash_string(k);
}
```

Implementing our Example

*You are the new produce manager of the local grocery store.
You want to use a dictionary to track your fruit inventory.*

- We can now implement the inventory application that uses hash dictionaries

- ✓ new dictionary
- ✓ insert A = ("apple", 20)
- ✓ insert B = ("banana", 10)
- ✓ insert C = ("pumpkin", 50)
- ✓ look up "apple"
- ✓ look up "lime"
- ✓ insert D = ("banana", 20)

- We store this code in a file called **produce-main.c0**

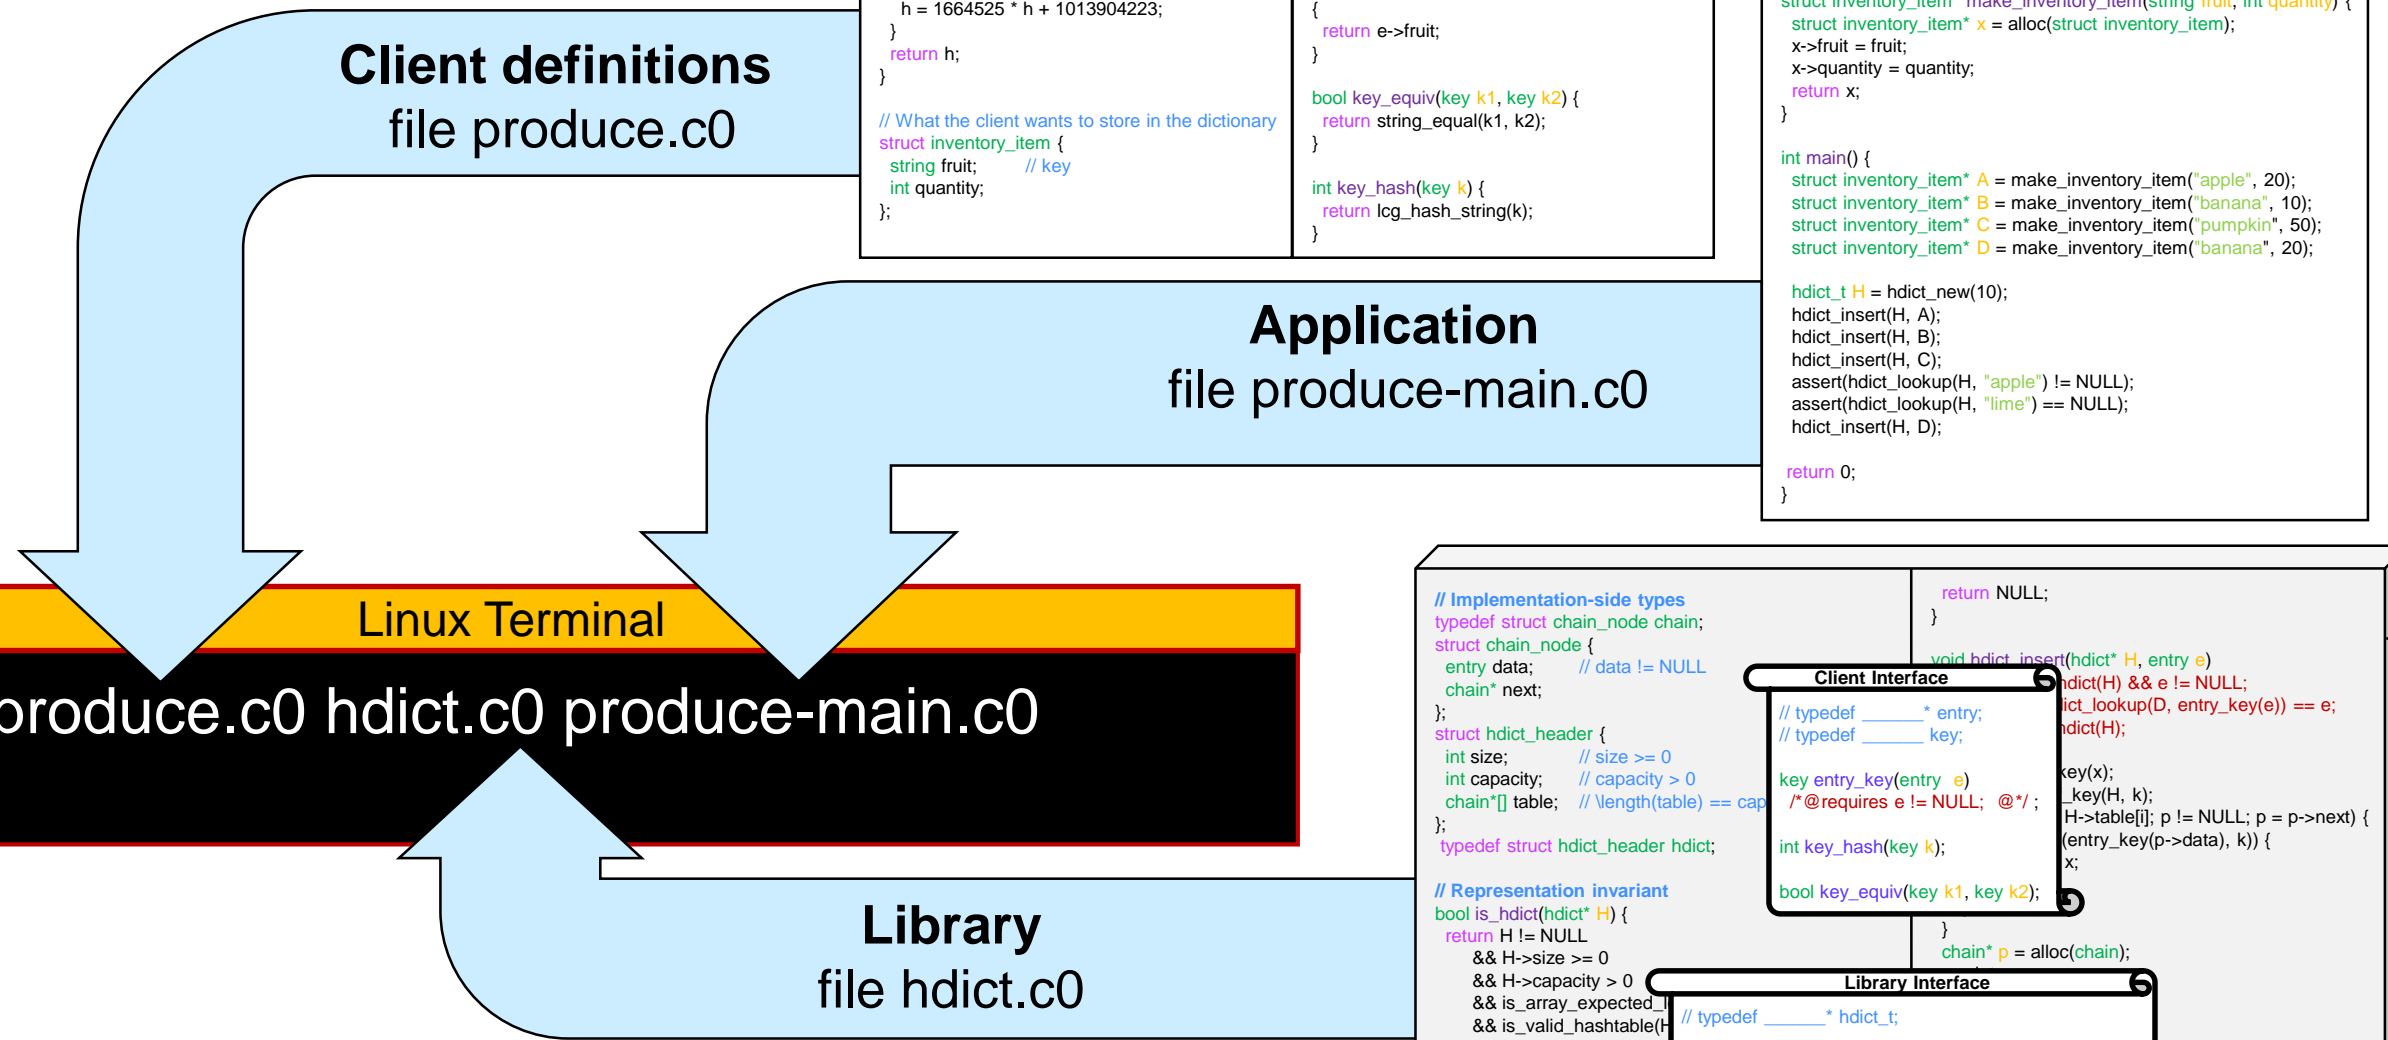
Client application file

```
struct inventory_item* make_inventory_item(string fruit, int quantity) {  
    struct inventory_item* x = alloc(struct inventory_item);  
    x->fruit = fruit;  
    x->quantity = quantity;  
    return x;  
}
```

Function that creates
inventory items

```
int main() {  
    struct inventory_item* A = make_inventory_item("apple", 20);  
    struct inventory_item* B = make_inventory_item("banana", 10);  
    struct inventory_item* C = make_inventory_item("pumpkin", 50);  
    struct inventory_item* D = make_inventory_item("banana", 20);  
  
    hdict_t H = hdict_new(10);  
    hdict_insert(H, A);  
    hdict_insert(H, B);  
    hdict_insert(H, C);  
    assert(hdict_lookup(H, "apple") != NULL);  
    assert(hdict_lookup(H, "lime") == NULL);  
    hdict_insert(H, D);  
  
    return 0;  
}
```

Compilation



- The definition file comes **before** the library
 - the library needs the definitions it supplies
- The library comes **before** the application file
 - the application needs the functionalities it provides

Compilation

Linux Terminal

```
# cc -d produce.c hdict.c produce-main.c
```

- *The definition file comes **before** the library*
 - *the library needs the definitions it supplies*
- *The library comes **before** the application file*
 - *the application needs the functionalities it provides*
- The client must split the application code into two files
 - This leads to an unnatural compilation pattern
 - We would like to compile the hash dictionary library just the way we compile a stack library



We will address this shortly

Hash Sets

Towards an Interface

What about Sets?

- **keys = entries**
 - these are the elements of the set
 - a single type **elem** replaces **key** and **entry**
- **lookup can simply return true or false**
 - this now checks set membership
 - return type is **bool**
 - no need to signal “not found” in a special way
 - **elem** does not have to be a pointer type

- A **set** can be understood as a special case of a dictionary
 - keys = entries
 - these are the elements of the set
 - **lookup** can simply return true or false
 - this now checks set membership
- A set implemented as a hash dictionary is called a **hash set**

The Hash Set Interface

Client Interface

```
// typedef _____ elem;  
  
int elem_hash(elem k);  
  
bool elem_equiv(elem k1, elem k2);
```

Library Interface

```
// typedef _____* hset_t;  
  
hset_t hset_new(int capacity)  
/*@requires capacity > 0;          @*/  
/*@ensures \result != NULL;        @*/;  
  
bool hset_contains(hset_t S, elem e)  
/*@requires S != NULL;             @*/;  
  
void hset_insert(hset_t S, elem e)  
/*@requires S != NULL;             @*/  
/*@ensures hset_contains(S, e);    @*/;
```

- A single type **elem** replaces **key** and **entry**
 - it does not need to be a pointer
- lookup checks membership
 - renamed **hset_contains**
 - it returns a **bool**
- Everything else remains the same

*The implementation
is left as exercise*