

Generic Hash Dictionaries

Hash Dictionaries so far

The Hash Dictionary Library

```

// Implementation-side types
typedef struct chain_node chain;
struct chain_node {
    entry data;          // data != NULL
    chain* next;
};
struct hdict_header {
    int size;             // size >= 0
    int capacity;         // capacity > 0
    chain*[] table;       // \length(table) == capacity
};
typedef struct hdict_header hdict;

// Representation invariant
bool is_hdict (hdict* H) {
    return H != NULL
        && H->size >= 0
        && H->capacity > 0
        && is_array_expected_length(H->table, H->capacity)
        && is_valid_hashtable(H);
}

// Implementation of interface functions
int index_of_key(hdict* H, key k)
/*@requires is_hdict(H);
  @ensures 0 <= \result && \result < H->capacity;
  {
    return abs(key_hash(k) % H->capacity);
  }

entry hdict_lookup(hdict* H, key k)
/*@requires is_hdict(H);
  @ensures \result == NULL
    || key_equiv(entry_key(\result), k);
  {
    int i = index_of_key(H, k);
    for (chain* p = H->table[i]; p != NULL; p = p->next)
        if (key_equiv(entry_key(p->data), k))
            return p->data;
    return NULL;
  }

void hdict_insert(hdict* H, entry e)
/*@requires is_hdict(H) && e != NULL;
  @ensures hdict_lookup(H, entry_key(e)) == e;
  @ensures is_hdict(H);
  {
    key k = entry_key(e);
    int i = index_of_key(H, k);
    for (chain* p = H->table[i]; p != NULL; p = p->next) {
        if (key_equiv(entry_key(p->data), k)) {
            p->data = e;
            return;
        }
    }
    chain* p = alloc(chain);
    p->data = e;
    p->next = H->table[i];
    H->table[i] = p;
    (H->size)++;
  }

hdict* hdict_new(int capacity)
/*@requires capacity > 0;
  @ensures is_hdict(\result);
  {
    hdict* H = alloc(hdict);
    H->size = 0;
    H->capacity = capacity;
    H->table = alloc_array(chain*, capacity);
    return H;
  }

// Client type
typedef hdict* hdict_t;

```

Implementation

Client Interface

```

// typedef _____* entry;
// typedef _____ key;

key entry_key(entry e)
/* @requires e != NULL; @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);

```

Library Interface

```



// typedef _____* hdict_t;

hdict_t hdict_new(int capacity)
/* @requires capacity > 0; @*/
/* @ensures \result != NULL; @*/;

entry hdict_lookup(hdict_t D, key k)
/* @requires D != NULL; @*/
/* @ensures \result != NULL
  || key_equiv(entry_key(\result), k); @*/;

void hdict_insert(hdict_t D, entry e)
/* @requires D != NULL && e != NULL; @*/
/* @ensures hdict_lookup(D, entry_key(e)) == e; @*/;

```



Is this Library Generic?

- Generic data structures

- work the same way no matter the type of their data

- Dictionaries are intrinsically generic

- they map keys to entries
- they work the same for any type of key and entry

- *Hash* dictionaries should be generic

- they abstract key manipulations into client functions
- `key_entry`, `key_hash` and `key_equiv`

- Generic libraries

- a single implementation that
 - lets the clients choose the type of their data
 - allows multiple instances of the data structure with different data types in the same application

Is this Library Generic?

- A single implementation that
 - lets the clients choose the type of their data
- **Yes!**
- the client interface mandates that the client define the types **key** and **entry**

Client Interface

```
// typedef _____* entry;
// typedef _____ key;

key entry_key(entry e)
/*@requires e != NULL; @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);
```

- the client does so in the client definition file
- Let's try it out to be sure

Client definition file

```
}
return h;
}

// What the client wants to store in the dictionary
struct inventory_item {
    string fruit;    // key
    int quantity;
};

/****** Fulfilling the library interface *****/
typedef struct inventory_item* entry;
typedef string key;

key entry_key(entry e)
/*@requires e != NULL;
{
    return e->fruit;
}

bool key_equiv(key k1, key k2) {
    return string_equal(k1, k2);
}

int key_hash(key k) {
    return lcg_hash_string(k);
}
```

Is this Library Generic?

- A single implementation that
 - lets the clients to chose the type of their data



Compiles and run the code

Linux Terminal

```
# cc0 -dx produce.c0 hdict.c0 produce-main.c0
All produce tests passed!
0
# cc0 -dx lib/*.c0 words.c0 hdict.c0 words-main.c0
All word count tests passed!
0
```

Another client application
that uses the hash dictionary
to count the occurrences of
each word in a file

Is this Library Generic?

- A single implementation that
 - lets the clients to chose the type of their data ✓
 - allows multiple instances of the data structure with different data types in the same application

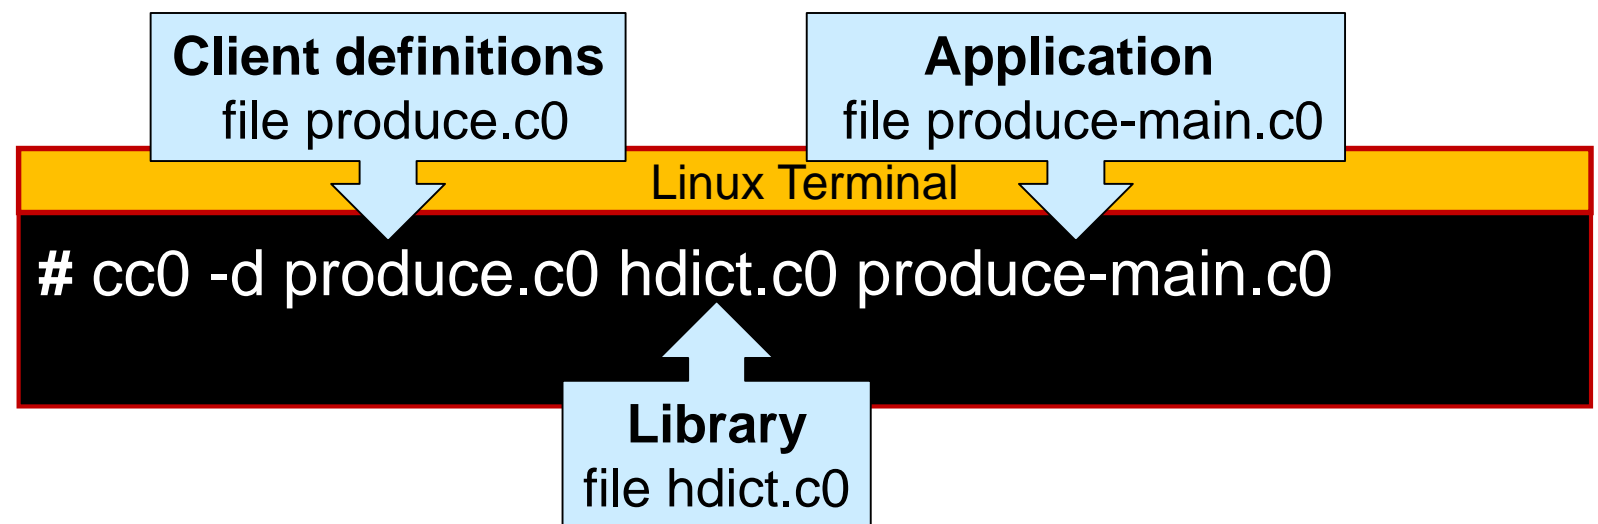
```
Linux Terminal
# cc0 -dx produce.c0 lib/*.c0 words.c0 hdict.c0 combined-main.c0
words.c0:29.1-29.30:error:type name 'entry' defined more than once
previous definition at produce.c0:28.1-28.38
typedef struct wcount* entry;
~~~~~
Compilation failed
```

➤ there can be at most one definition of the types **key** and **entry**



Is this Library Generic?

- *A single implementation that*
 - *lets the clients choose the type of their data* ✓
 - *allows multiple instances of the data structure with different data types in the same application* ✗
- This approach also forced clients to split their application code into two files



- This is an unnatural compilation pattern
 - We would like to compile the hash dictionary library just the way we compile a stack library

Making this Library Generic

- With the stack library, setting the element type to `void*` solved both problems

```
typedef void* elem;
```

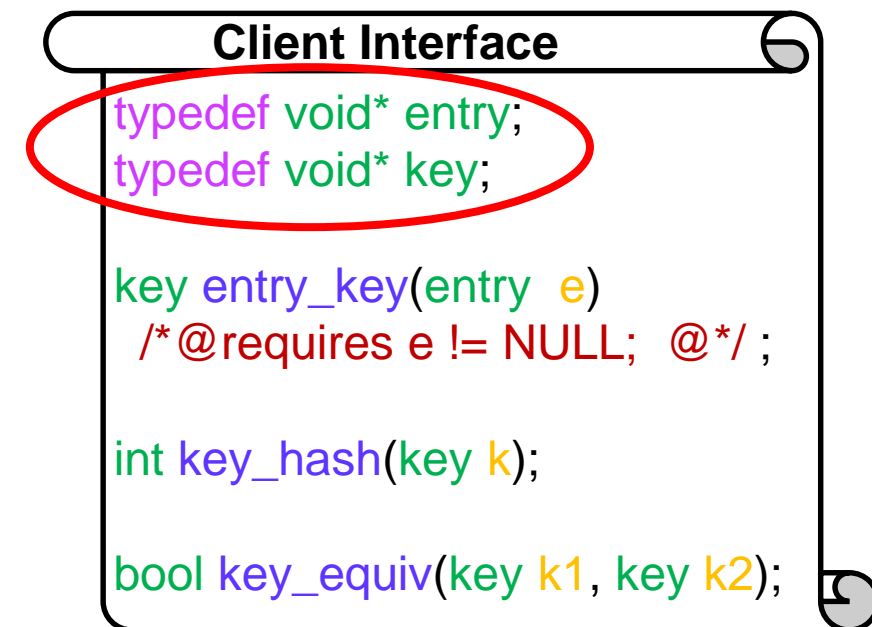

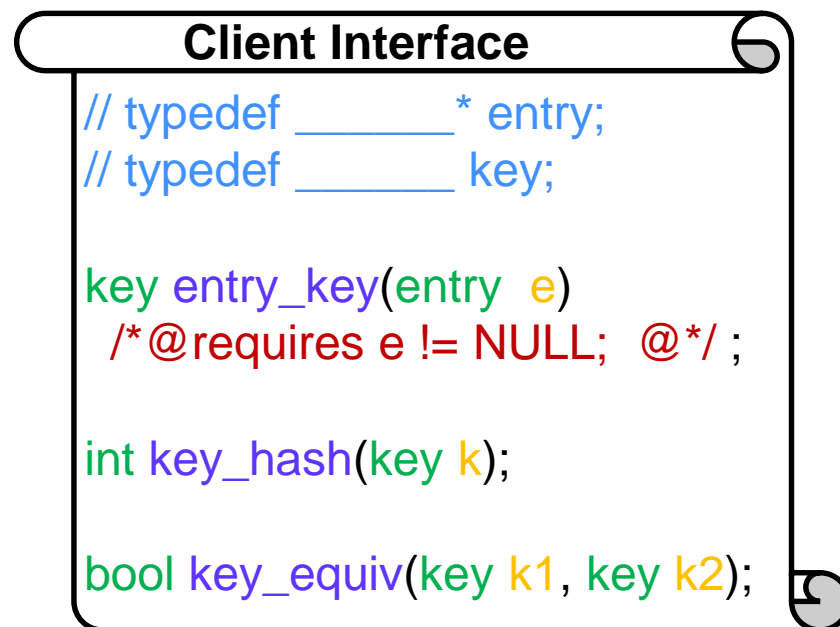
➤ This was now C1 code

- Let's do the same with the hash dictionary library

void*
to the Rescue

Upgrading the Library

- The **only** changes we need to make to the library are defining **key** and **entry** as **void***



- This only affects the client interface

We should not add NULL-check on keys since the client could have

- chosen key to be a pointer type
- chosen NULL to be a valid key

That's it!

Upgrading the Client Definitions

Client Interface


```
typedef void* entry;
typedef void* key;

key entry_key(entry e)
/* @requires e != NULL; @ */;

int key_hash(key k);

bool key_equiv(key k1, key k2);
```

- The client does not need to define **key** and **entry**
 - the library defines both to **void***



- For the client
 - an entry is still a **struct inventory_item***
 - a key is now a **string*** It's got to be a pointer
 - "lime" must now live in a cell in allocated memory to be used as a key
 - NULL does not correspond to any valid key

```

typedef void* entry;
typedef void* key;

key entry_key(entry e)
/*@requires e != NULL; @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);

```


Upgrading the Client Definitions

- *For the client*
 - an entry is still a *struct inventory_item**
 - a key is now a *string**
- So,
 - every value of type *entry* must have tag *struct inventory_item**
 - every value of type *key* must have tag *string**

```

int key_hash(key k) {
    return lcg_hash_string(k);
}

```



```

int key_hash(key k)
/*@requires k != NULL && \hastag(string*, k);
{
    ...
}

```

NULL is not a
valid fruit name

Every *key* is in
reality a *string**


Upgrading the Client Definitions

Client Interface

```
typedef void* entry;  
typedef void* key;  
  
key entry_key(entry e)  
/*@requires e != NULL; @*/;  
  
int key_hash(key k);  
  
bool key_equiv(key k1, key k2);
```

- For the client, a key is now a *string**
- Before using a value of type *key*, we need to
 - cast it to *string**
 - dereference the result to a *string*

```
int key_hash(key k) {  
    return lcg_hash_string(k);  
}
```



```
int key_hash(key k)  
/*@requires k != NULL && \hastag(string*, k);  
{  
    return lcg_hash_string(*(string*)k);  
}
```

lcg_hash_string takes
a *string* as input

A *void** cast
to a *string** and
dereferenced to a *string*

```

typedef void* entry;
typedef void* key;

key entry_key(entry e)
/*@requires e != NULL; @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);

```


Upgrading the Client Definitions

- *For the client*
 - an entry is still a *struct inventory_item**
 - a key is now a *string**
 - *"lime"* must now live in a cell in allocated memory to be used as a key
 - *NULL* does not correspond to any valid key
- When extracting a key from an entry, we must put it in a cell in allocated memory

```

key entry_key(entry e)
/*@requires e != NULL;
{
    return e->fruit;
}

```



```

key entry_key(entry e)
/*@requires e != NULL && \hastag(struct inventory_item*, e);
/*@ensures \result != NULL && \hastag(string*, \result);
{
    struct inventory_item* E = (struct inventory_item*)e;
    string* K = alloc(string);
    *K = E->fruit;
    return (key)K;
}

```

Auxiliary variables can help make those casts more readable

```

typedef void* entry;
typedef void* key;

key entry_key(entry e)
/*@requires e != NULL; @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);

```


Upgrading the Client Definitions

- When extracting a key from an entry, we must put it in a cell in allocated memory
 - add a helper function that turns a `string` to a `string*` for even better readability

```

key entry_key(entry e)
/*@requires e != NULL;
{
    return e->fruit;
}

```



```

string* to_string_ptr(string s)
/*@ensures \result != NULL;
{
    string* s_ptr = alloc(string);
    *s_ptr = s;
    return s_ptr;
}

key entry_key(entry e)
/*@requires e != NULL && \hastag(struct inventory_item*, e);
/*@ensures \result != NULL && \hastag(string*, \result);
{
    struct inventory_item* E = (struct inventory_item*)e;
    string* K = to_string_ptr(E->fruit);
    return (key)K;
}

```

We could write it in one line as
`return (key) to_string_ptr(((struct inventory_item*)e)->fruit);`
 but that's unreadable

Upgrading the Client Definitions

Client Interface

```
typedef void* entry;
typedef void* key;

key entry_key(entry e)
/*@requires e != NULL; @*/;

int key_hash(key k);

bool key_equiv(key k1, key k2);
```

```
// What the client wants
struct inventory_item {
    string fruit;    // key
    int quantity;
};
```


```
/****** Fulfilling the library interface *****/
typedef struct inventory_item* entry;
typedef string key;
```

```
key entry_key(entry e)
/*@requires e != NULL;
{
    return e->fruit;
}
```

```
bool key_equiv(key k1, key k2) {
    return string_equal(k1, k2);
}
```

```
int key_hash(key k) {
    return lcg_hash_string(k);
}
```

Similar to `key_hash`



```
// What the client wants to store
struct inventory_item {
    string fruit;    // key
    int quantity;
};
```

```
/****** Fulfilling the library interface *****/
```

```
key entry_key(entry e)
/*@requires e != NULL && \hastag(struct inventory_item*, e);
/*@ensures \result != NULL && \hastag(string*, \result);
{
    struct inventory_item* E = (struct inventory_item*)e;
    string* K = to_string_ptr(E->fruit);
    return (key)K;
}
```


```
bool key_equiv(key k1, key k2)
/*@requires k1 != NULL && \hastag(string*, k1);
/*@requires k2 != NULL && \hastag(string*, k2);
{
    return string_equal(*(string*)k1, *(string*)k2);
}
```

```
int key_hash(key k)
/*@requires k != NULL && \hastag(string*, k);
{
    return lcg_hash_string(*(string*)k);
}
```

Upgrading the Client Application

- Cast entries and keys before calling the library operations
- Turn values of type `string` to `string*` before using them as keys

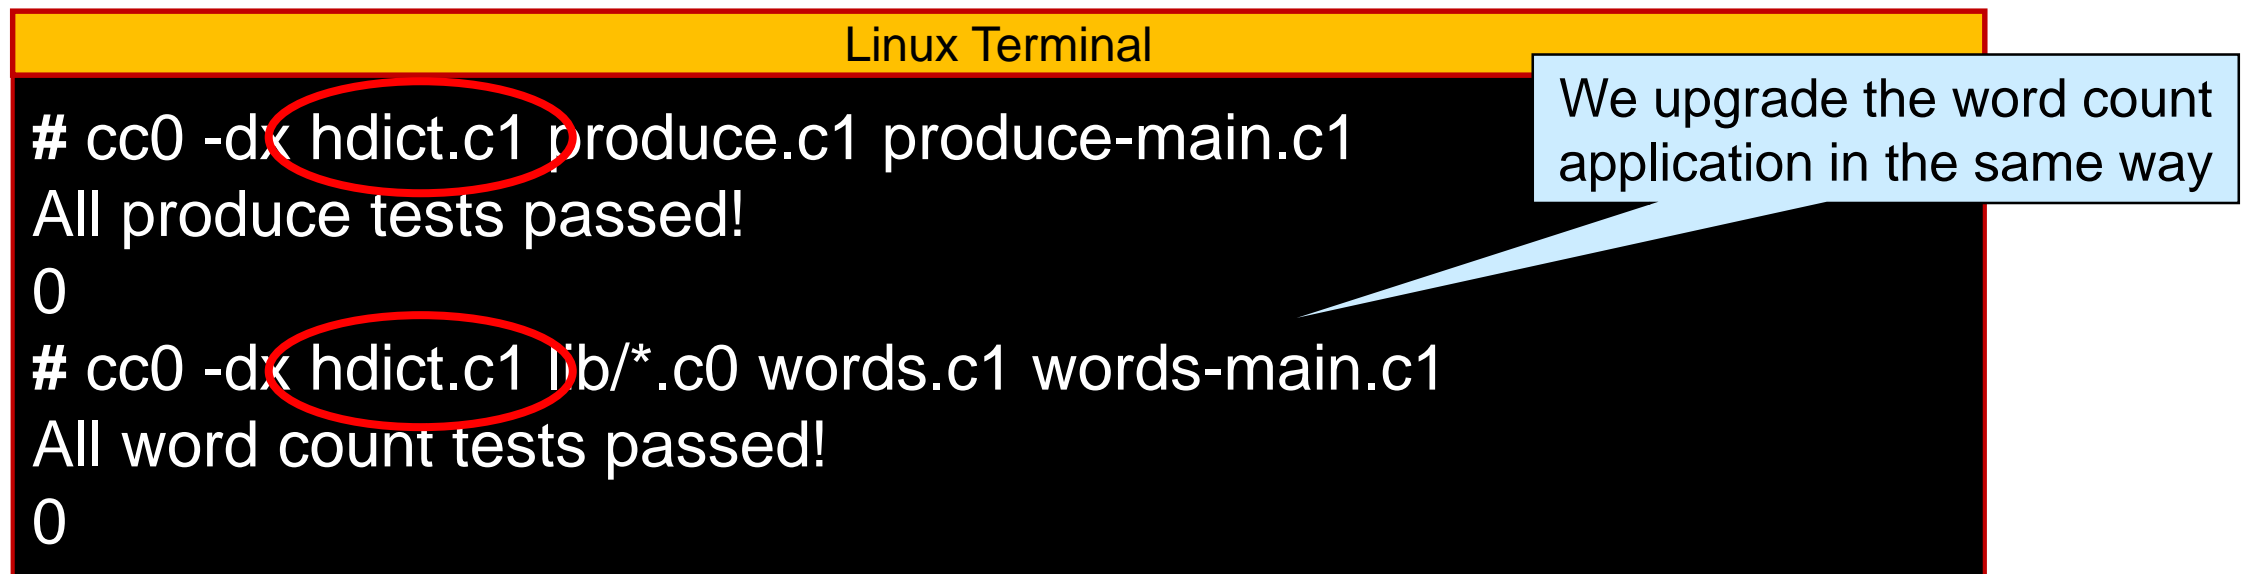
```
int main () {  
    struct inventory_item* A = make_inventory_item("apple", 20);  
    struct inventory_item* B = make_inventory_item("banana", 10);  
    struct inventory_item* C = make_inventory_item("pumpkin", 50);  
    struct inventory_item* D = make_inventory_item("banana", 20);  
  
    hdict_t H = hdict_new(10);  
    hdict_insert(H, A);  
    hdict_insert(H, B);  
    hdict_insert(H, C);  
    assert(hdict_lookup(H, "apple") != NULL);  
    assert(hdict_lookup(H, "lime") == NULL);  
    hdict_insert(H, D);  
  
    return 0;  
}
```



```
int main () {  
    struct inventory_item* A = make_inventory_item("apple", 20);  
    struct inventory_item* B = make_inventory_item("banana", 10);  
    struct inventory_item* C = make_inventory_item("pumpkin", 50);  
    struct inventory_item* D = make_inventory_item("banana", 20);  
  
    hdict_t H = hdict_new(10);  
    hdict_insert(H, (entry)A);  
    hdict_insert(H, (entry)B);  
    hdict_insert(H, (entry)C);  
    assert(hdict_lookup(H, (key)to_string_ptr("apple")) != NULL);  
    assert(hdict_lookup(H, (key)to_string_ptr("lime")) == NULL);  
    hdict_insert(H, (entry)D);  
  
    return 0;  
}
```

Generic Hash Dictionaries

- Let's try it on our examples



```
Linux Terminal
# cc0 -dx hdict.c1 produce.c1 produce-main.c1
All produce tests passed!
0
# cc0 -dx hdict.c1 lib/*.c0 words.c1 words-main.c1
All word count tests passed!
0
```

We upgrade the word count application in the same way

- we can now put the library **before** all client files
- this means we could merge produce.c1 and produce-main.c1 into a single file
 - same thing for word.c1 and words-main.c1

Generic Hash Dictionaries

- Let's try it on our examples

```
Linux Terminal
# cc0 -dx hdict.c1 produce.c1 lib/*.c0 words.c1 combined-main.c1
words.c1:38.1-45.2:error:function 'entry_key' defined more than once
previous definition at produce.c1:31.1-38.2
key entry_key(entry x) ... }
~~~~~
Compilation failed
```




- This still doesn't work!
 - both produce.c1 and words.c1 define `entry_key`
 - and `key_equiv` and `key_hash`
 - This is not allowed in C0/C1
 - Even if it were, the library wouldn't know which version to use with what data

Function Pointers to the Rescue

Renaming the Client Functions

- We avoid having duplicate client definition function names by renaming them
 - `key_hash` to `key_hash_produce`
 - etc
- and similarly for the word count application



- But how to tell the library which function to use?

This is all we need to do in the client definition file

```
// What the client wants to store
struct inventory_item {
    string fruit;    // key
    int quantity;
};

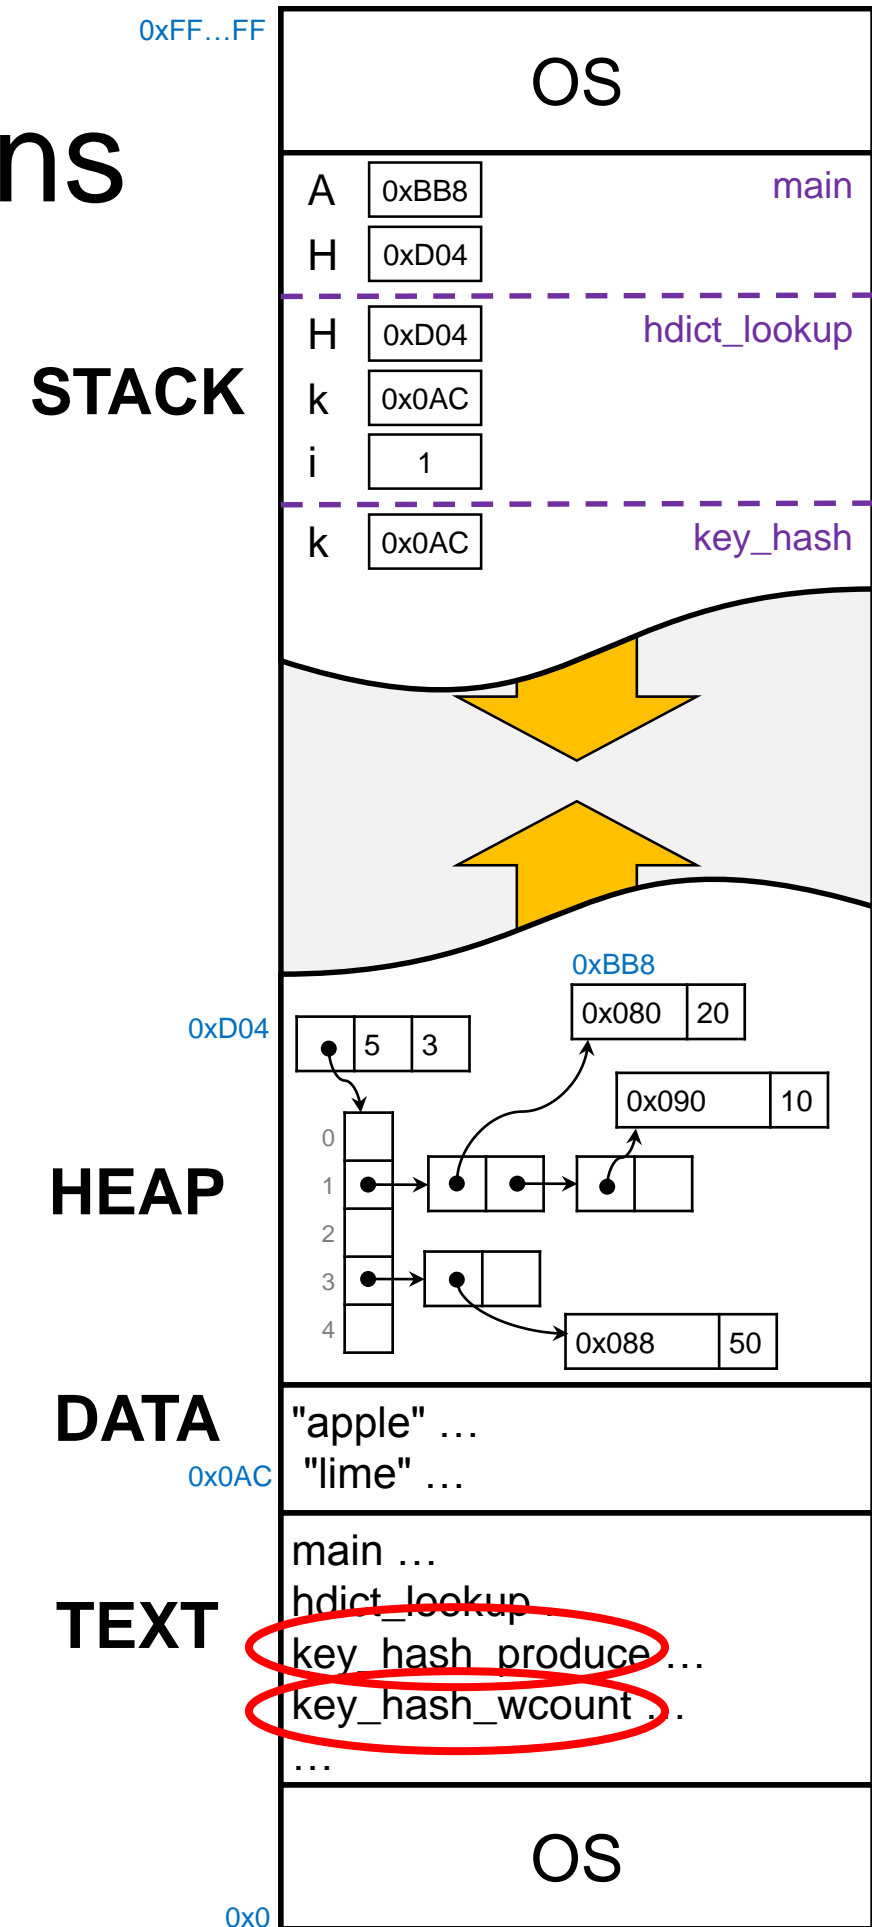
/***** Fulfilling the library interface *****/
key_entry_key_produce(entry e)
//@requires e != NULL && \hastag(struct inventory_item*, e);
//@ensures \result != NULL && \hastag(string*, \result);
{
    struct inventory_item* E = (struct inventory_item*)e;
    string* K = to_string_ptr(E->fruit);
    return (key)K;
}

bool key_equiv_produce(key k1, key k2)
//@requires k1 != NULL && \hastag(string*, k1);
//@requires k2 != NULL && \hastag(string*, k2);
{
    return string_equal(*(string*)k1, *(string*)k2);
}

int key_hash_produce(key k)
//@requires k != NULL && \hastag(string*, k);
{
    return lcg_hash_string(*(string*)k);
}
```

Accessing the Right Functions

- During execution, functions live in the TEXT segment of memory
 - & allows us to obtain their address and pass it around as a function pointer
 - we can call a function through a pointer to it
- **Idea:** make pointers to the appropriate client function available to the library
 - *but how to do so?*



Accessing the Right Functions – I

- One option is to pass the right client functions to the library functions that use them

- So,

```
entry A = make_inventory_item("apple", 20);  
hdict_insert(H, A);
```

becomes

```
entry A = make_inventory_item("apple", 20);  
hdict_insert(H, A, &entry_key_produce,  
             &key_hash_produce,  
             &key_equiv_produce);
```


- Then do this for every use of `hdict_insert` and `hdict_lookup`
 - We will make mistakes
 - This is also a poor way of thinking about dictionaries
 - H is a dictionary where we want to store produce
 - every insertion and lookup on H will need these client functions
 - not others



Accessing the Right Functions – II

- A better option is to pass the right client functions when we *create* a dictionary
 - in `hdict_new`

```
hdict_t H = hdict_new(10);
```



```
hdict_t H = hdict_new(H, &entry_key_produce  
&key_hash_produce,  
&key_equiv_produce);
```

This is all we need to do
in the client application file

- `hdict_new` needs to store the client functions in H itself
 - we need to modify the internal representation
 - but first we need to give types to the client functions

Client Function Types

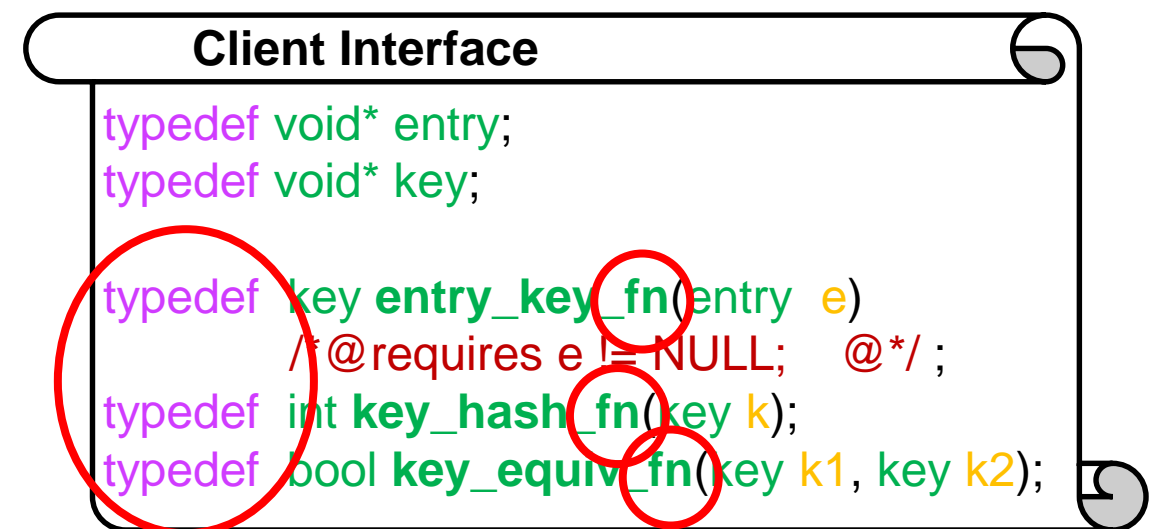
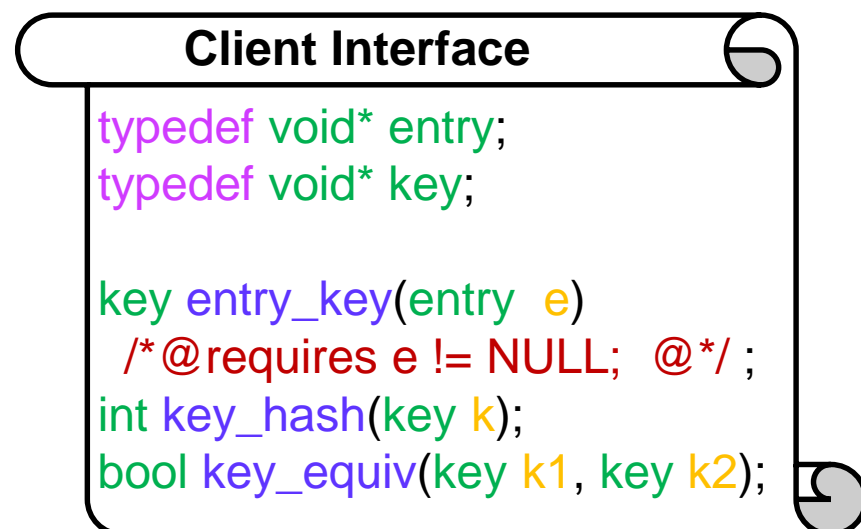
- We need to define types for the client functions

- `entry_key_fn`
- `key_hash_fn`
- `key_equiv_fn`

since

- `hdict_new` takes them as arguments
- we will store them in the concrete implementation type


- These definitions go in the client interface



Upgrading the Concrete Type

- We store the client definitions in the data structure itself
 - extend `struct hdict_header` with three additional fields

```
struct hdict_header {  
    int size;           // size >= 0  
    int capacity;       // capacity > 0  
    chain*[] table;     // \length(table) == capacity  
};  
typedef struct hdict_header hdict;
```



```
struct hdict_header {  
    int size;           // size >= 0  
    int capacity;       // capacity > 0  
    chain*[] table;     // \length(table) == capacity  
    entry_key_fn* key;  // != NULL  
    key_hash_fn* hash;  // != NULL  
    key_equiv_fn* equiv; // != NULL  
};  
typedef struct hdict_header hdict;
```


- Storing both data and functions in a struct is a fundamental concept in **object-oriented programming**
 - these structs are called **objects**
 - the functions are called **methods**

There is a lot more to object-oriented programming however

Upgrading the Representation Invariant

- A valid **hdict** cannot have NULL in the added fields

```
bool is_hdict (hdict* H) {  
    return H != NULL  
        && H->size >= 0  
        && H->capacity > 0  
        && is_array_expected_length(H->table, H->capacity)  
        && is_valid_hashtable(H);  
}
```




```
bool is_hdict (hdict* H) {  
    return H != NULL  
        && H->size >= 0  
        && H->capacity > 0  
        && is_array_expected_length(H->table, H->capacity)  
        && H->key   != NULL  
        && H->hash  != NULL  
        && H->equiv != NULL  
        && is_valid_hashtable(H);  
}
```

Upgrading `hdict_new`

- `hdict_new`

- takes the client functions as inputs
- expects them to be non-NULL
- stores them in the added fields of the concrete type

```
hdict* hdict_new(int capacity)
//@requires capacity > 0;
//@ensures is_hdict(\result);
{
    hdict* H = alloc(hdict);
    H->size = 0;
    H->capacity = capacity;
    H->table = alloc_array(chain*, capacity);
    return H;
}
```




```
hdict* hdict_new(int capacity,
    entry_key_fn* entry_key,
    key_hash_fn* hash,
    key_equiv_fn* equiv)
//@requires capacity > 0;
//@requires entry_key != NULL && hash != NULL && equiv != NULL;
//@ensures is_hdict(\result);
{
    hdict* H = alloc(hdict);
    H->size = 0;
    H->capacity = capacity;
    H->table = alloc_array(chain*, capacity);
    H->key = entry_key;
    H->hash = hash;
    H->equiv = equiv;
    return H;
}
```

Calling the Client Functions

- Whenever we need a client function, we call the function pointer in the data structure
 - E.g., `key_hash`
 - Here, `H->hash`
- For example,

```
int index_of_key(hdict* H, key k)
//@requires is_hdict(H);
//@ensures 0 <= \result && \result < H->capacity;
{
    return abs(key_hash(k) % H->capacity);
}
```




```
int index_of_key(hdict* H, key k)
//@requires is_hdict(H);
//@ensures 0 <= \result && \result < H->capacity;
{
    return abs((*H->hash)(k) % H->capacity);
}
```

This is the same as
`(*(H->hash))(...)`

Upgrading `hdict_lookup`

- Proceed in the same way
 - change client function calls to function pointer calls

```
entry hdict_lookup(hdict* H, key k)
//@requires is_hdict(H);
//@ensures \result == NULL
           || key_equiv(entry_key(\result), k);
{
  int i = index_of_key(H, k);
  for (chain* p = H->table[i]; p != NULL; p = p->next) {
    if (key_equiv(entry_key(p->data), k))
      return p->data;
  }
  return NULL;
}
```




```
entry hdict_lookup(hdict* H, key k)
//@requires is_hdict(H);
//@ensures \result == NULL
           || key_equiv(entry_key(\result), k);
{
  int i = index_of_key(H, k);
  for (chain* p = H->table[i]; p != NULL; p = p->next) {
    if ((*H->equiv)((*H->key)(p->data), k))
      return p->data;
  }
  return NULL;
}
```

- The function pointer syntax is hard to read
 - factor it out in helper functions similar to `index_of_key`

Upgrading `hdict_lookup`

- *The function pointer syntax is hard to read*
 - *factor it out in helper functions similar to `index_of_key`*

```
entry hdict_lookup(hdict* H, key k)
//@requires is_hdict(H);
//@ensures \result == NULL
           || key_equiv(entry_key(\result), k);
{
    int i = index_of_key(H, k);
    for (chain* p = H->table[i]; p != NULL; p = p->next) {
        if (key_equiv(entry_key(p->data), k))
            return p->data;
    }
    return NULL;
}
```



```
int index_of_key(hdict* H, key k)
//@requires is_hdict(H);
//@ensures 0 <= \result && \result < H->capacity;
{
    return abs((*H->hash)(k) % H->capacity);
}

key entry_key(hdict* H, entry x)
//@requires is_hdict(H) && x != NULL;
{
    return (*H->key)(x);
}

bool key_equiv(hdict* H, key k1, key k2)
//@requires is_hdict(H);
{
    return (*H->equiv)(k1, k2);
}

entry hdict_lookup(hdict* H, key k)
//@requires is_hdict(H);
//@ensures \result == NULL
           || key_equiv(entry_key(\result), k);
{
    int i = index_of_key(H, k);
    for (chain* p = H->table[i]; p != NULL; p = p->next) {
        if (key_equiv(H, entry_key(H, p->data), k))
            return p->data;
    }
    return NULL;
}
```

- `hdict_insert` is similar

Updating the Library Interface

- The client interface does not contain `entry_key` and `key_equiv` any more

- We cannot call them in the library interface

```
Library Interface
entry hdict_lookup(hdict_t D, key k)
/* @requires D != NULL;                               @*/
/* @ensures \result != NULL
   || key_equiv(entry_key(\result, k); @*/;
```



- Using the fields of the implementation type would violate the interface

```
Library Interface
entry hdict_lookup(hdict_t D, key k)
/* @requires D != NULL;                               @*/
/* @ensures \result != NULL
   || (*H->equiv)((*H->key)(\result, k); @*/;
```



- We must give up on this refined postcondition

```
Library Interface
entry hdict_lookup(hdict_t D, key k)
/* @requires D != NULL;                               @*/;
```



The Hash Dictionary Library

```
// Implementation-side types
typedef struct chain_node chain;
struct chain_node {
    entry data;          // data != NULL
    chain* next;
};

struct hdict_header {
    int size;             // size >= 0
    int capacity;         // capacity > 0
    chain*[] table;       // \length(table) == capacity
    entry_key_fn* key;     // != NULL
    key_hash_fn* hash;    // != NULL
    key_equiv_fn* equiv;  // != NULL
};

typedef struct hdict_header hdict;

// Representation invariant
bool is_hdict(hdict* H) {
    return H != NULL
        && H->size >= 0
        && H->capacity > 0
        && is_array_expected_length(H->table, H->capacity)
        && H->key != NULL
        && H->hash != NULL
        && H->equiv != NULL
        && is_valid_hashtable(H);
}

// Implementation of interface functions
int index_of_key(hdict* H, key k)
/*@requires is_hdict(H);
@ensures 0 <= \result && \result < H->capacity;
{
    return abs((*H->hash)(k) % H->capacity);
}

key entry_key(hdict* H, entry x)
/*@requires is_hdict(H) && x != NULL;
{
    return (*H->key)(x);
}
```

```
bool key_equiv(hdict* H, key k1, key k2)
/*@requires is_hdict(H);
{
    return (*H->equiv)(k1, k2);
}

entry hdict_lookup(hdict* H, key k)
/*@requires is_hdict(H);
@ensures \result == NULL
        || key_equiv(entry_key(\result), k);
{
    int i = index_of_key(H, k);
    for (chain* p = H->table[i]; p != NULL; p = p->next) {
        if (key_equiv(H, entry_key(H, p->data), k))
            return p->data;
    }
    return NULL;
}

void hdict_insert(hdict* H, entry e) // left as exercise

hdict* hdict_new(int capacity,
    entry_key_fn* entry_key, key_hash_fn* hash,
    key_equiv_fn* equiv)
/*@requires capacity > 0 && entry_key != NULL;
@requires hash != NULL && equiv != NULL;
@ensures is_hdict(\result);
{
    hdict* H = alloc(hdict);
    H->size = 0;
    H->capacity = capacity;
    H->table = alloc_array(chain*, capacity);
    H->key = entry_key;
    H->hash = hash;
    H->equiv = equiv;
    return H;
}

// Client type
typedef hdict* hdict_t;
```

Implementation

Client Interface

```
typedef void* entry;
typedef void* key;

typedef key entry_key_fn(entry e)
    /* @requires e != NULL; @*/;
typedef int key_hash_fn(key k);
typedef bool key_equiv_fn(key k1, key k2);
```

Library Interface

```
// typedef _____* hdict_t;

hdict_t hdict_new(int capacity,
    entry_key_fn* entry_key,
    key_hash_fn* hash,
    key_equiv_fn* equiv)
    /* @requires capacity > 0 && entry_key != NULL; @*/
    /* @requires hash != NULL && equiv != NULL; @*/
    /* @ensures \result != NULL; @*/;

entry hdict_lookup(hdict_t D, key k)
    /* @requires D != NULL; @*/;

void hdict_insert(hdict_t D, entry e)
    /* @requires D != NULL && e != NULL; @*/;
```

Is it Generic?

Linux Terminal

```
# cc0 -dx hdict.c1 produce.c1 lib/*.c0 words.c1 combined-main.c1  
All word count tests passed!  
All produce tests passed!  
0
```

- Yes!



Harmony

Experimenting with Client Definitions

- Now that we have an easy way to specify which client definition functions to use, we can test a few things
 - Let's consider alternatives versions of `key_hash` and `key_equiv`
 - they look at the length of the key

This is a bad hash function.
Let's use it anyway
for simplicity

```
int key_hash_produce(key k)
//@requires k != NULL && \hastag(string*, k);
{
    return lcg_hash_string(*(string*)k);
}
```

```
int key_hash_produce_alt(key k)
//@requires k != NULL && \hastag(string*, k);
{
    return string_length(*(string*)k);
}
```

```
bool key_equiv_produce(key k1, key k2)
//@requires k1 != NULL && \hastag(string*, k1);
//@requires k2 != NULL && \hastag(string*, k2);
{
    return string_equal(*(string*)k1, *(string*)k2);
}
```

```
bool key_equiv_produce_alt(key k1, key k2)
//@requires k1 != NULL && \hastag(string*, k1);
//@requires k2 != NULL && \hastag(string*, k2);
{
    return string_length(*(string*)k1) == string_length(*(string*)k2);
}
```

This are our original
functions

This are our alternative
functions

Mixing and Matching

- `key_hash_produce` and `key_equiv_produce` are meant to be used together

```
hdict_t H hdict_new(H, &entry_key_produce,  
                    &key_hash_produce, &key_equiv_produce);
```

- Same for `key_hash_produce_alt` and `key_equiv_produce_alt`

```
hdict_t H hdict_new(H, &entry_key_produce,  
                    &key_hash_produce_alt, &key_equiv_produce_alt);
```

- But what if we mix and match them?

Mixing and Matching

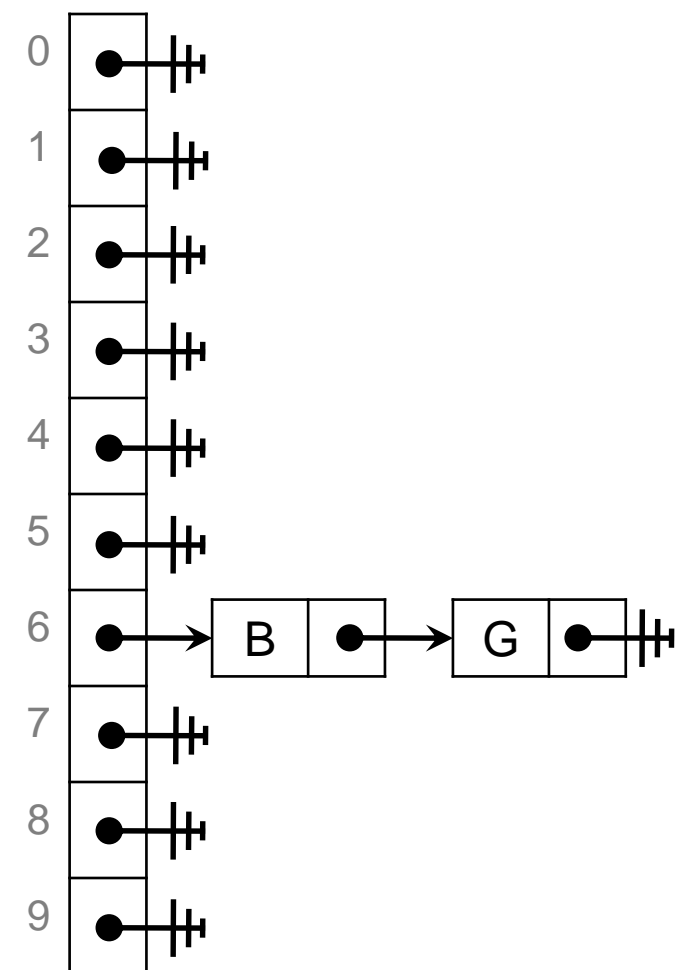
- But what if we mix and match them?
 - `key_hash_produce_alt` with `key_equiv_produce`

```
hdict_t H hdict_new(H, &entry_key_produce,  
                    &key_hash_produce_alt, &key_equiv_produce);
```

- Let's use the dictionary

- ✓ new dictionary
- ✓ insert B = ("banana", 10)
- ✓ insert G = ("grapes", 30)

- `key_hash_produce_alt` returns 6 on both
 - both "banana" and "grapes" have length 6
- both end up in the same bucket
 - but `key_hash_produce` would have sent them in different buckets
- This is not as **efficient** as using `key_hash_produce`
 - the dictionary works correctly, but not is not as fast



Mixing and Matching

- But what if we mix and match them?

- `key_hash_produce` with `key_equiv_produce_alt`

```
hdict_t H hdict_new(H, &entry_key_produce,  
                    &key_hash_produce, &key_equiv_produce_alt);
```

That's the other way around

- Let's use the dictionary


✓ new dictionary
✓ insert B = ("banana", 10)
look up "grapes"

`key_hash_produce`
returns 7 on "banana"

- ❑ `key_hash_produce` returns 2 on "grapes"
- ❑ there is nothing in bucket 2
- ❑ lookup "grapes" returns NULL

➤ **This is incorrect!**

- "grapes" and "banana" have length 6
- `key_equiv_produce_alt` treats them as equal
- ❑ What *look up "grapes"* asks is *find an entry whose key has length 6*
 - B fits the bill, but it is not found



X

Harmony

- The key hash and equivalence functions are in **harmony** when equivalent entries have the same hash value
 - `key_hash_produce` and `key_equiv_produce_alt` are **not** in harmony
 - ❑ `"banana"` and `"grapes"` are equal according to `key_equiv_produce_alt`
 - ❑ but, according to `key_hash_produce`
`"banana"` hashes to index 7 while `"grapes"` hashes to index 2
 - the other combinations are in harmony
- When they are not in harmony, the hash dictionary does not work correctly
 - it may return the wrong answer
- Harmony does not guarantee efficiency